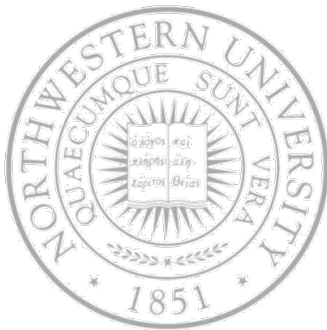


Architectural Support for Operating Systems



Today

- Computer system overview

Next time

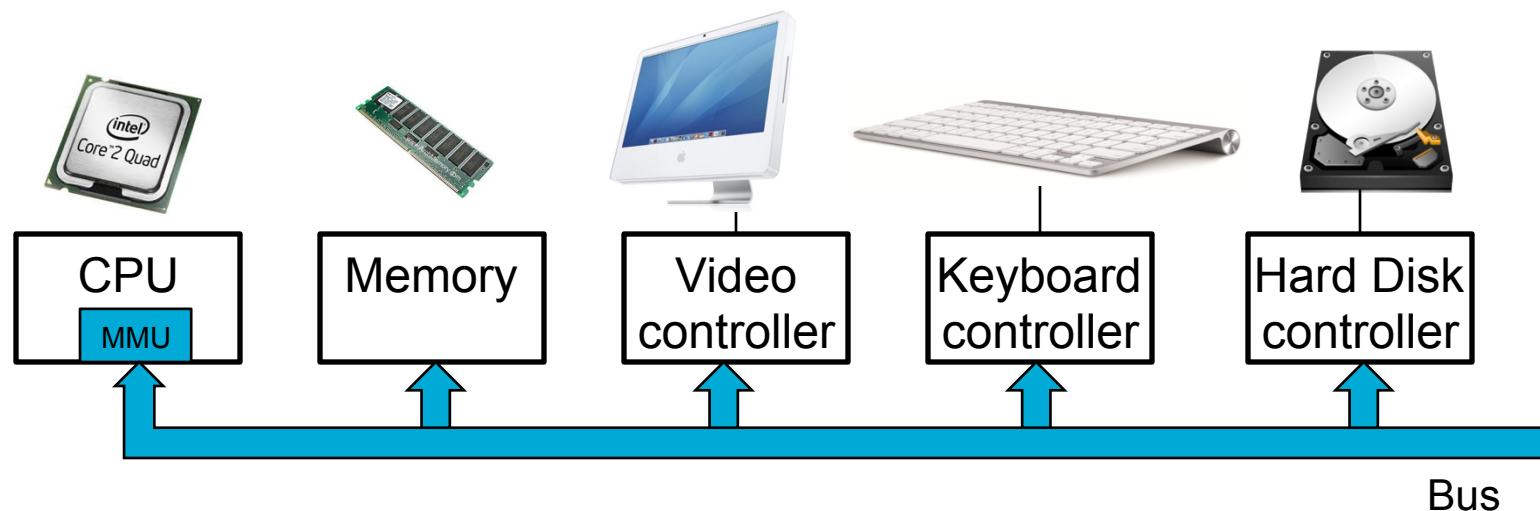
- OS components & structure

Announcements and reminders

- In case you have missed anything ...
- Project 1 is out!

Computer architecture and OS

- OS is intimately tied to the hardware it runs on
 - OS design is impacted by it
 - OS needs result on new architectural features
- Abstract model of a simple computer




Processor

- The brain with a basic operation cycle
 - Fetch next instruction
 - Decode it to determine type & operands
 - Execute it
- ... and a specific set of instructions
 - E.g. combine ops (ADD), control flow, data movement
 - Architecture specific - Pentium != SPARC
- Since memory access is slow ... registers
 - General regs to hold variables & temp. results
 - Special regs such as Program Counter, Stack Pointer




Processor ...

- This model is overly simplistic: pipeline architectures, superscalar, ...





- Multithreading/Hyperthreading and multicore



Memory

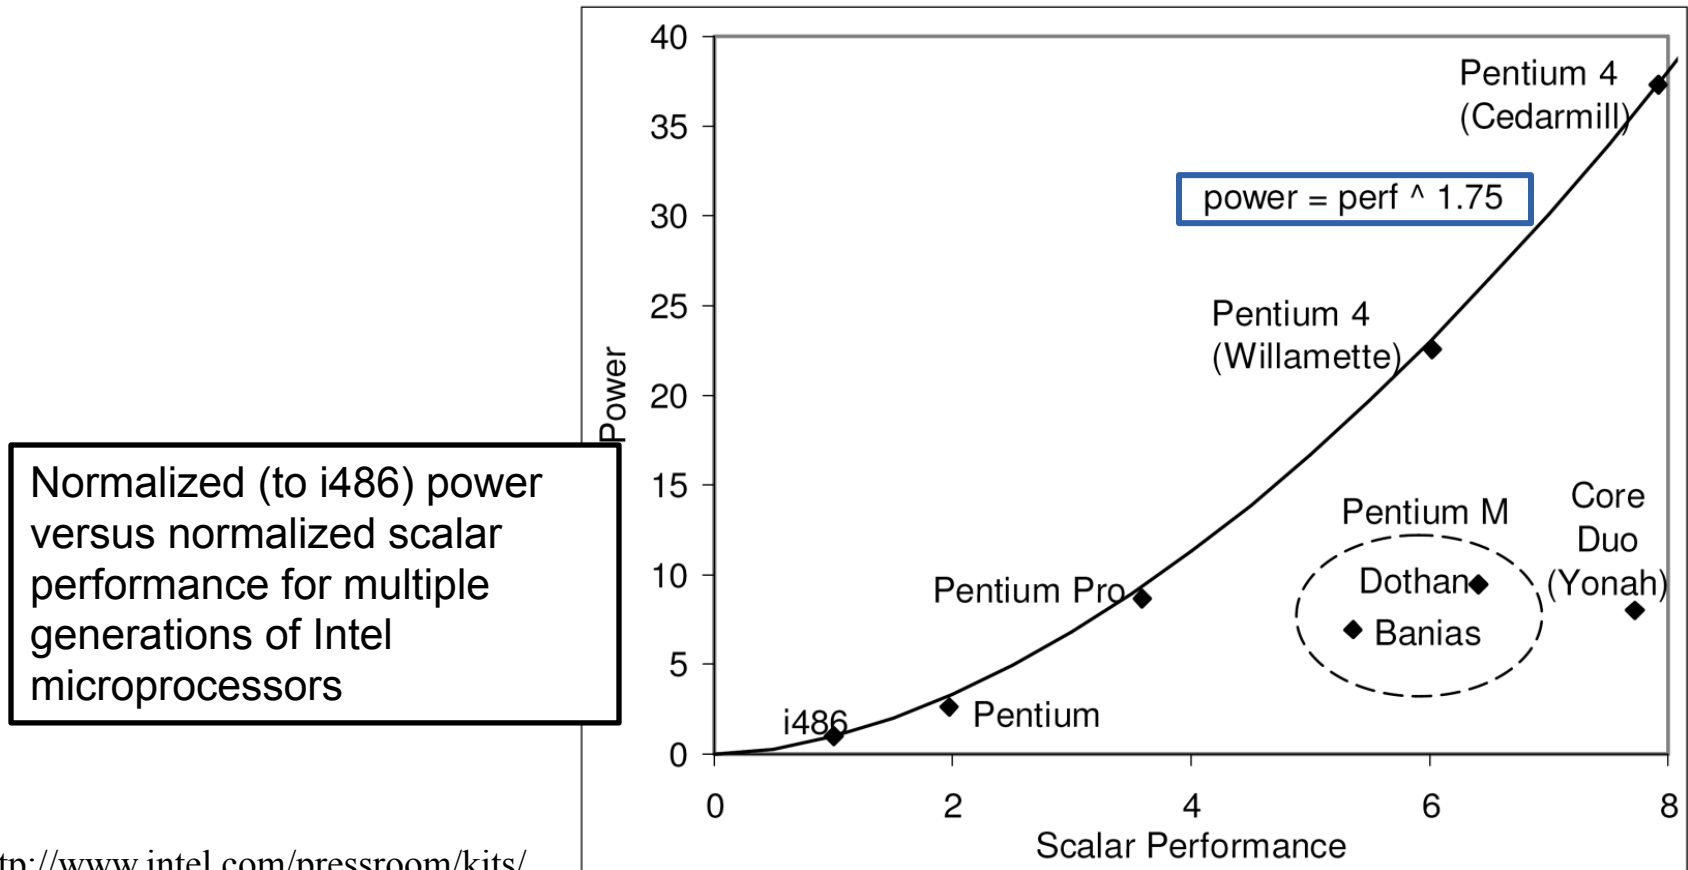
- Ideal – fast, large, cheap and persistent
- Real – storage hierarchy
 - Registers
 - Internal to the CPU & just as fast
 - 64x64 in a 64 bit machine
 - Cache
 - Split into cache lines
 - If word needs is in cache, get in ~2 cycles
 - Main memory
 - Hard disk
 - Magnetic tape
 - *Coherency?*

First core-based memory: IBM 405
Alphabetical Accounting Machine



Architectural trends impact OS design ...

- Processing power
 - Doubling every 18 months (100x per decade)
 - but power is a serious issue




*<http://www.intel.com/pressroom/kits/core2duo/pdf/epi-trends-final2.pdf>

Architectural trends impact OS design ...

- Primary memory capacity

- Same and for the same reason

1980	64KB	\$405.00 (\$6,480/MB)
1990	8MB	\$851.00 (\$106/MB)
2000	64MB	\$99.89 (\$1.56/MB)
2009	4GB	\$39.99 (\$0.010/MB)
2014	4GB	\$29.99 (\$0.007/MB)*




*<http://www.jcmit.com/memoryprice.htm>


Architectural trends impact OS design ...

- Disk capacity

- Double every 12 months (1000x per decade)



1961 IBM 1301
~26MB ~\$115,500 (~890k in 2013)



Disk pack
Weight: ~10lbs
Capacity: 2MB



2014 WD 3TB My
Book Essential ~ \$120

1961 ~ \$4,440/MB
➔
2014 ~ \$0.00004/MB

Architectural trends impact OS design

- Solid state storage (SSD)
 - 10-100k random IOs per second
 - 800 MB/s transfer rates
 - Costly, but quickly riding Moore's law
- | | | |
|------|-------------------|-------|
| 2011 | Crucial 512GB SSD | \$750 |
| 2012 | ... | \$400 |
| 2014 | ... | \$210 |




Crucial 512 GB m4
2.5-Inch Solid State
Drive SATA 6Gb/s



*Hard to imagine with hard drives –
“The SSD revolution”, arstechnica
June 25, 2012*

Architectural trends impact OS design ...

- Gap between CPU and I/O speeds



Relative speed of key components –
an unbalanced system
FAWN project @ CMU

Architectural trends impact OS design

- Optical bandwidth today
 - Doubling every 9 months (Butter's law)
 - 50% improvement each year for home users (Nielsen's law)
 - Factor of 10,000 every decade
 - 10x as fast as disk capacity!
 - 100x as fast as processor performance!
- What are some of the implications of these trends?
 - E.g.: from mainframes to desktops to cloud computing

And now a short break ...

Loop



... and OS needs shape the architecture

- Arch support can simplify/complicate OS tasks
 - E.g., early PC OS (DOS, MacOS) lacked support for virtual memory, partly because HW lacked key features
- Features built primarily to support OS's:
 - Protected modes of execution (kernel vs. user)
 - Protected instructions
 - System calls (and software interrupts)
 - Memory protection
 - I/O control operations
 - Timer (clock) operation
 - Interrupts and exceptions
 - Synchronization instructions

Consider timesharing

- Multiprogramming & timesharing are useful
 - Multiprogramming – “using different parts of the hardware at the same time for different tasks”
 - Timesharing – “several persons making use of the computer at the same time”
- but
 - How to protect programs from each other & kernel from all?
 - How to handle relocation? OS may need to run a particular program at different times from several locations

OS protection

- Some instructions are restricted to the OS
 - e.g. Directly access I/O devices, manipulate memory state management
- How does the CPU know if a protected instructions should be executed?
 - Architecture must support 2+ mode of operation
 - Mode is set by status bit in a protected register (PSW)
 - User programs execute in user mode, OS in kernel mode
- Protected instructions can only execute in kernel mode

Crossing protection boundaries

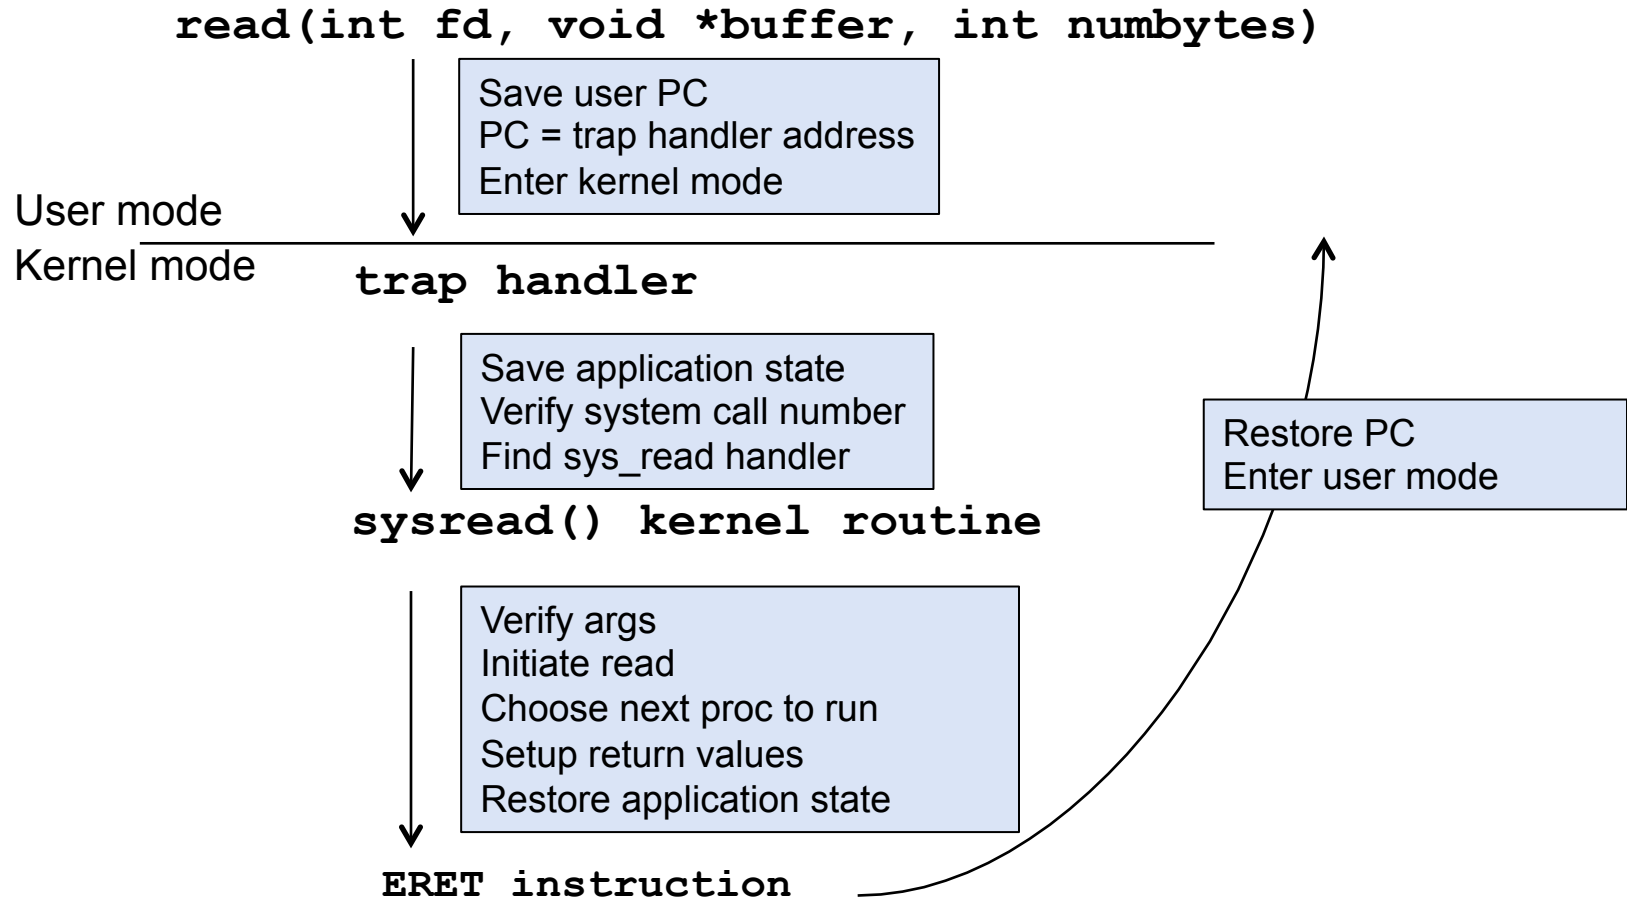
- How can apps do something privileged?
 - e.g. How do you save a file if you can't do I/O?
- User programs must call an OS procedure
 - Ask the OS to do it
 - OS defines a set of **system calls**
 - User-mode program makes a system call
 - How does the user to kernel-mode transition happen?

Crossing protection boundaries

- The system call ...
 - Causes an exception which vector to a kernel handler
 - Passes a parameter indicating which syscall is
 - Saves caller's state so it can be restored
 - *What would happen if the kernel didn't save state?*
 - OS must verify caller's parameters
 - *Why should it do that?*
 - Must be a way to go back to user once done
 - A special instruction sets PC to the return address and the execution mode to user
- *A bit like a regular subroutine call, right?*

Crossing protection boundaries

- A system call




Exception handling and protection

- All entries to the OS use the same mechanism
 - Acquiring privileged mode and branching to trap handler are inseparable
- Interrupts, exceptions and traps
 - Interrupt: asynchronous, caused by an external hw event
 - Exception: synchronous; unexpected, automatically generated trap (coerced rather than requested); divide by zero
 - Trap: synchronous; programmer initiated, expected transfer of control to a special handler
- Privileged instructions and resources are the basis for most everything: Memory protection, I/O ...

Next issue – Memory relocation

- Simplest model – base + limit
 - Base (start) of program + limit registers
 - Used by both CDC 6600 (first supercomputer) and Intel 8088, hear of the IBM PC
 - Changing program means changing base+limit
 - Solves relocation and protection (check on the CDC 6600)
 - Cost 2 registers + cycle time incr
- More sophisticated alternatives
 - 2 base and 2 limit registers for text & data; allow sharing program text
 - Paging, segmentation, virtual memory



If \geq base and $<$ base + limit, OK
else trap to OS with error

OS needs shape the architecture – I/O

- I/O Device
 - Device + Controller (simpler I/F to OS; think SCSI)
 - Read sector x from disk y → (disk, cylinder, sector, head), ...
- How does the kernel start an I/O?
 - Special I/O instructions
 - Memory-mapped I/O
- How does it notice when the I/O is done?
 - Polling – *are we done yet?*
 - Interrupts – *let me know when you are done?*
- How does it exchange data with the I/O device?
 - Programmed I/O
 - Direct Memory Access (DMA)

OS control flow

- OSs are event driven
 - Once booted, all entry to kernel happens as result of an event (e.g. signal by an interrupt), which
 - Immediately stops current execution
 - Changes to kernel mode, event handler is called
- Kernel defines handlers per event type
 - Specific types are defined by the architecture
 - e.g. timer event, I/O interrupt, system call trap

Interrupts and exceptions

- Three main types of events: interrupts & exceptions
 - Exceptions/traps caused by SW executing instructions
 - E.g., a page fault
 - E.g., an attempted write to a read-only page
 - An expected exception is a “trap”, unexpected is a “fault”
 - Interrupts caused by HW devices
 - E.g., device finishes I/O
 - E.g., timer fires

Timers

- How can the OS retain control when a program gets stuck in an infinite loop?
 - Use a HW timer that generates a periodic interrupt
 - Before it transfers to a user program, the OS loads the timer with a time to interrupt (how long?)
 - When time's up, interrupt transfers control back to OS
 - OS pick a program to schedule next (which one?)
- Should the timer be privileged?
 - For reading or for writing?

Synchronization

- Issues with interrupts

- Many at a time, executing code can interfere with interrupted code
- OS must be able to synchronize concurrent processes

- Synchronization

- Guarantee that short instruction sequences (e.g. read-modify-write) execute atomically
- Two methods
 - Turn off interrupts, execute sequence, re-enable interrupts
 - Have special, complex atomic instructions – test-and-set

Management of concurrency & asynchronous events is the biggest difference between systems-level & traditional app programming

Summary

- This is far from over – new architectural features are still being introduced
 - Support for virtual machine monitors
 - Hardware transaction support
 - Support for security
 - ...
- Transistors are free so Intel/AMD/... need to find applications that require new hardware that you would want to buy ...