


OS Concepts and structure



Today

- OS services
- OS interface to programmers/users
- OS components & interconnects
- Structuring OSs

Next time

- Processes

OS Views

- Vantage points
 - OS as the services it provides
 - To users and applications
 - OS as its components and interactions
- OS provides a number of services
 - To users via a command interpreter/shell or GUI
 - To application programs via system calls
 - Some services are for convenience
 - Program execution, I/O operation, file system management, communication
 - Some to ensure efficient operation
 - Resource allocation, accounting, protection and security

Command interpreter (shell) & GUI

- Command interpreter

- Handle (interpret and execute) user commands
- Could be part of the OS: MS DOS, Apple II
- Could be just a special program: UNIX, Win XP
 - In this way, multiple shells are possible
- The command interpreter could
 - Implement all commands
 - Simply understand what program to invoke and how (UNIX)

- GUI

- Friendlier (desktop), if sometimes limiting
- Xerox PARK Alto >> Apple >> Windows >> Linux

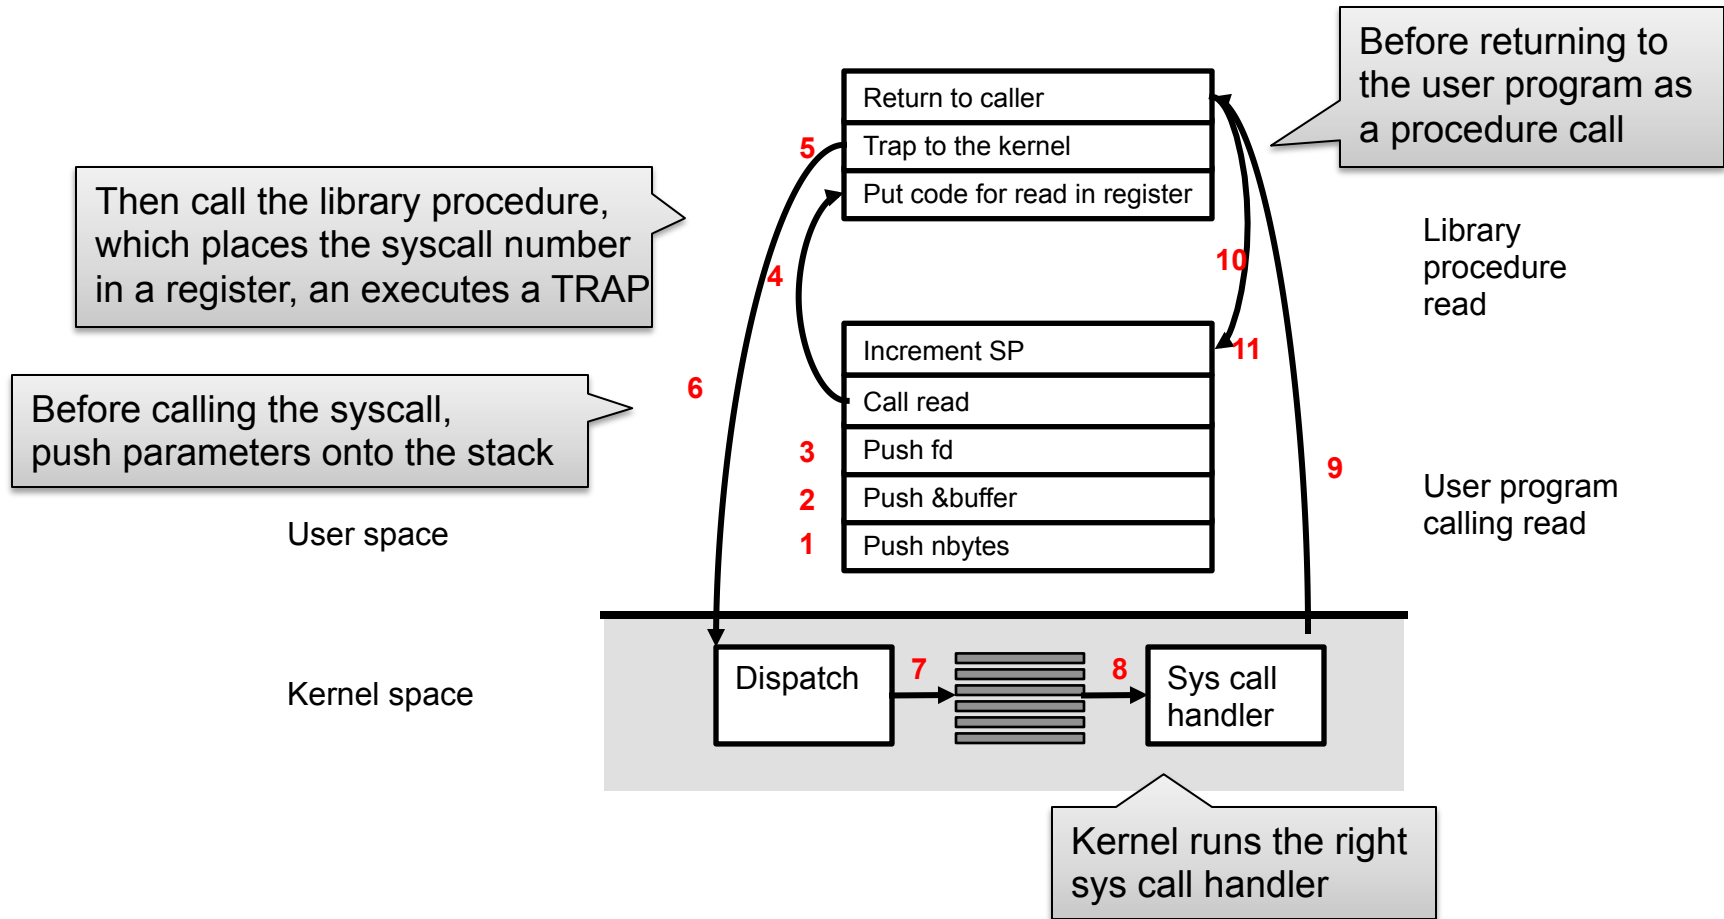
System calls

- Low-level interface to services for applications
- Higher-level requests get translated into sequence of system calls
- Writing `cp` – copy source to destination
 - Get file names
 - Open source
 - Create destination
 - Loop
 - Read from source
 - Copy to destination
 - Close destination
 - Report completion
 - Terminate

System calls

- The steps in making a read system call

```
count = read(fd, buffer, nbytes);
```



Major OS components & abstractions

- Processes
- Memory
- I/O
- Secondary storage
- File systems
- Protection
- Accounting
- Shells & GUI
- Networking

Processes

- A program in execution
 - Address space
 - Set of registers
 - Threads and processes – for now consider each process to have a single thread (we'll change that later)
 - Linux: `ps -auwwx` to list all processes
- To get a better sense of it
 - What data do you need to (re-) start a suspended process?
 - Where do you keep this data?
 - What is the process abstraction interface offered by the OS?
 - Create, delete, suspend, resume & clone a process
 - Inter-process communication & synchronization

Memory management

- Main memory – the directly accessed storage for CPU
 - Programs must be stored in memory to execute
 - Memory access is fast (e.g., 60 ns to load/store)
 - but memory doesn't survive power failures
- OS must
 - Allocate memory space to processes (explicitly and implicitly)
 - Decide how to allocate to each process
 - Deallocate space when needed by rest of system
 - Maintain mappings from physical to virtual memory
 - Decide when to remove a process from memory

I/O

- A big chunk of the OS kernel deals with I/O
 - Hundreds of thousands of lines in NT
- The OS provides a standard interface between programs & devices
 - file system (disk), sockets (network), frame buffer (video)
- Device drivers are the routines that interact with specific device types
 - Encapsulates device-specific knowledge
 - e.g., how to initialize a device, request I/O, handle errors
 - Examples: SCSI device drivers, Ethernet card drivers, video card drivers, sound card drivers, ...

Secondary storage

- Secondary storage (disk, tape) is persistent memory
 - Often magnetic media, survives power failures (hopefully)
- Routines that interact with disks are typically at a very low level in the OS
 - Used by many components (file system, VM, ...)
 - Handle scheduling of disk operations, head movement, error handling, and often management of space on disks
- Usually independent of file system
 - Although there may be cooperation
 - File system knowledge of device details can help optimize performance
 - e.g., place related files close together on disk

File systems

- Storage devices are hard to work with
 - File system offers a convenient abstraction
 - Defines logical abstractions/objects like files & directories
 - As well as operations on these objects
- A file is the basic unit of long-term storage
- A directory is just a special kind of file
 - ... containing names of other files & metadata
- Interface
 - File/directory creation/deletion, manipulation, copy, lock
- Other higher level services: accounting & quotas, backup, indexing or search, versioning

Protection

- Protection is a general mechanism used throughout the OS
 - All resources must be protected
 - memory
 - processes
 - files
 - devices
 - ...
- Protection mechanisms help to detect and contain errors, as well as preventing malicious destruction

And now a short break ...

OS structure

- OS made of number of components
 - Process & memory management, file system, ...
 - and system programs
 - e.g., bootstrap code, the init program, ...
- Major design issue
 - How do we organize all this?
 - What are the modules, and where do they exist?
 - How do they interact?
- Massive software engineering
 - Design a large, complex program that:
 - performs well, is reliable, extensible, backwards compatible, ...

OS design & implementation

- *User goals and System goals*
 - User – convenient to use, easy to learn, reliable, safe, fast
 - System – easy to design, implement, & maintain, also flexible, reliable, error-free & efficient
- Affected by choice of hardware, type of system
- Clearly conflicting goals, no unique solution
- Some other issues complicating this
 - Size: Windows XP ~40G SLOC, RH 7.1 17G SLOC
 - Concurrency – multiple users and multiple devices
 - Potentially hostile users, but some users want to collaborate
 - Long expected lives & no clear ideas on future needs
 - Portability and support to thousands of device drivers
 - Backward compatibility


OS design & implementation

- A software engineering principle – separate policy & mechanism
 - Policy: What will be done?
 - Mechanism: How to do it?
 - Why do you care? Max flexibility, easier to change
- Implementation on high-level language
 - Early on – assembly (e.g. MS-DOS – 8088), later Algol (MCP), PL/1 (MULTICS), C (Unix, ...)
 - Advantages – faster to write, more compact, easier to maintain & debug, easier to port
 - Cost – Size, speed?, but who cares?!

Early versions ... were written in assembly language, but during the summer of 1973, it was rewritten in C. The size of the new system is about one third greater than the old. ... much easier to understand and to modify but also includes many functional improvements ... we considered this increase in size quite acceptable.

D. Ritchie and K. Thompson, The UNIX time-sharing system, CACM 17(7), 1974

Monolithic design



- Major advantage
 - Cost of module interactions is low (procedure call)
- Disadvantages
 - Hard to understand
 - Hard to modify & maintain
 - Unreliable (no isolation between system modules)
- Alternative?
 - How to organize the OS in order to simplify design, implementation and maintenance?

Layering

- The traditional approach
 - Implement OS as a set of layers
 - Each layer shows an enhanced 'virtual mach' to layer above

Layer	Function
5	The operator
4	User programs
3	I/O management
2	Operator-process communication
1	Memory and drum management
0	Processor allocation and multiprogramming

Dijkstra's THE system

- Each layer can be tested and verified independently

Problems with layering

- Imposes hierarchical structure
 - but real systems have complex interactions
 - Strict layering isn't flexible enough
- Poor performance
 - Each layer crossing has an associated overhead
- Disjunction between model and reality
 - Systems modelled as layers, but not built that way

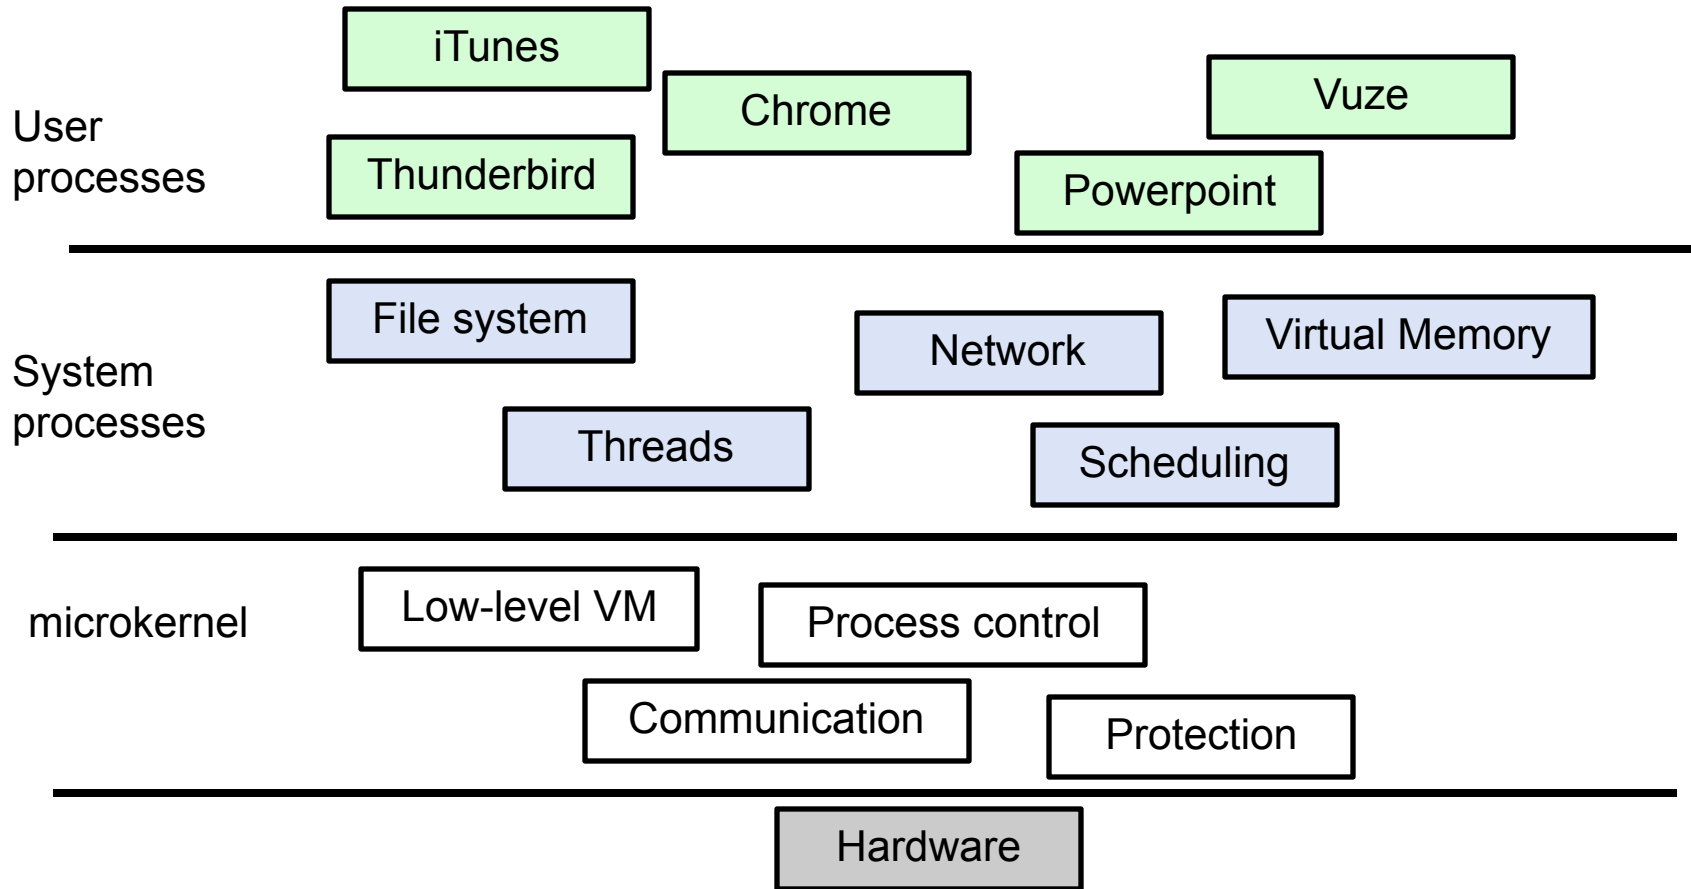
HAL – Hardware Abstraction Layer

- An example of layering in modern OSs
- Goal – to hide differences in hardware from most of the OS kernel
 - On a PC, you can consider it as the driver of the motherboard
 - BSD, Mac OS X, Windows NT, Linux, NetBSD all use a HAL either explicitly identified or not

Microkernels

- Popular in the late 80's, early 90's
 - Recent resurgence
- Goal
 - Minimize what goes in kernel
 - Organize rest of OS as user-level processes
- This results in
 - Better reliability (isolation between components)
 - Ease of extension and customization
 - Poor performance (user/kernel boundary crossings)
- First microkernel Hydra (CMU, 1970)
 - ... Mach (CMU), Chorus (UNIX-like), OS X (Apple), in some ways NT (Microsoft), L4 (Karlsruhe), MINIX 3, ...

Microkernel




Virtual machines

- Initial release of OS/360 were strictly batch but users wanted timesharing
 - IBM CP/CMS, later renamed VM/370 ('79)
- Note that timesharing systems provides
 1. Multiprogramming
 2. Extended (virtual) machine
- Essence of VM/370 – separate the two
 - Heart of the system (VMM) does multiprogramming & provides multiple exact copies of bare HW to next layer up
 - Each VM can run any OS

Virtual machines


- A resurgence in mid-90s, started with Rosenblum's work on Disco and VMWare
 - Nowadays ... Java VM, Xen, VritulaBox, Virtual Iron, VMLite, Simics, Parallels, Palacios, QEMU, ...
- What for?
 - Server consolidation – from different services in different lightly used machine to consolidation (administration cost)
 - Different applications for other OS in your desktop
 - Testing and debugging



Exokernels

- OS, typically securely multiplexes & *abstract* physical resources
- But no OS abstractions fits all!
- Exokernel
 - A minimal OS securely multiplexes resources
 - Library OSes implement higher-level abstractions

Secure binding – a protection mechanism that decouples authorization (done at binding) from use of a resource



Summary & preview

- Today
 - The mess under the carpet
 - Basic concepts in OS
 - OS design has been an evolutionary process
 - Structuring OS - a few alternatives, not a clear winner
- Next ...
 - Process – the central concept in OS
 - Process model and implementation
 - What it is, what it does and how it does it