


# Virtual Memory

---

## Today

- Handling bigger address spaces
- Speeding translation



# Considerations with page tables

## Two key issues with page tables

- Mapping must be fast
  - Done on every memory reference, at least 1 per instruction
- With large address spaces, page tables will be large
  - 32 bit addresses & 4KB page  $\rightarrow$  12 bit offset, 20 bit page #  $\sim$  1million PTE ( $2^{20}$ )
  - 64 bit addresses & 4KB page  $\rightarrow$  ...  $2^{52}$  pages  $\sim 4.5 \times 10^{15}$ !!!
- Simplest solutions
  - Page table in registers
    - Fast during execution, \$\$\$ & slow to context switch
  - Page table in memory & Page Table Base Register (PTBR)
    - Fast to context switch & cheap, but slow during execution

# Page table and page sizes

---

- Bigger pages
  - Smaller page tables
  - But more internal fragmentation
- Smaller pages
  - Less internal fragmentation
  - Less unused program in memory,
  - But ... more I/O time, getting page from disk ... most of the time goes into seek and rotational delay!
  - Larger page tables

# Page table and page sizes

- Average process size  $s$  bytes
- Page size  $p$  bytes
- Page entry size  $e$  bytes
- Approximate number of pages needed per process  $s/p$  occupying  $s/p * e$  bytes of page table space
- $overhead = se / p + p/2$



Internal  
fragmentation

- Taking first derivative respect to  $p$  and equating it to zero


$$-s e / p^2 + 1/2 = 0$$

$$p = \sqrt{2se}$$

- $s = 1\text{MB}$        $e = 8 \text{ bytes}$  → Optimal  $p = 4\text{KB}$


# Separate instruction & data spaces

- One address space – size limit
- Pioneered by PDP-11: 2 address spaces, Instruction and Data spaces
  - Double the space
  - Each with its own page table & paging algorithm



# A hybrid approach – pages and segments

- As MULTICS
  - Instead of a single page table, one per segment
  - The base register of the segment points to the base of the page table

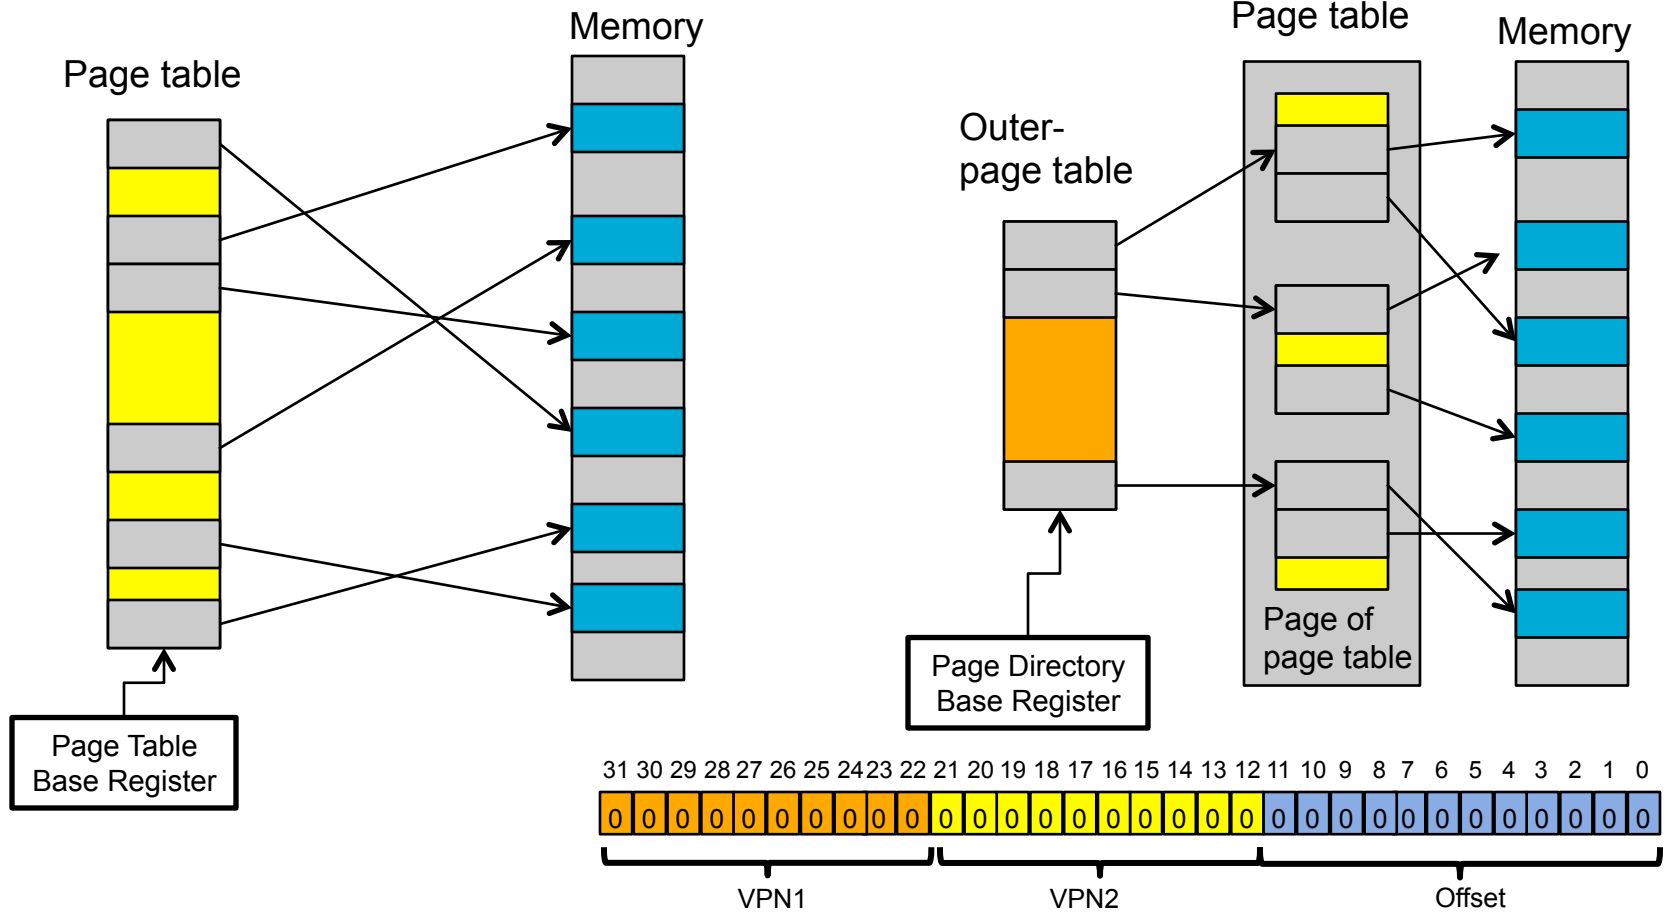



- But
  - Segmentation is not as flexible as we may want
  - Fragmentation is still an issue
  - Page tables can be of any size, so here we go again

# Hierarchical or multilevel page table

- Another approach – page the page table!
  - Same argument – you don't need the full page table
- Example
  - Virtual address (32-bit machine, 4KB page):  
Page # (20 bits) + Offset (12 bits)
  - Since page table is paged, page number is divided:  
Page number (10 bits) + Page offset in 2<sup>nd</sup> level (10 bits)
- Pros and cons
  - Allocate page table space as needed
  - If carefully done, each portion of the page table fits neatly within a page
  - More effort for translation
  - And a bit more complex

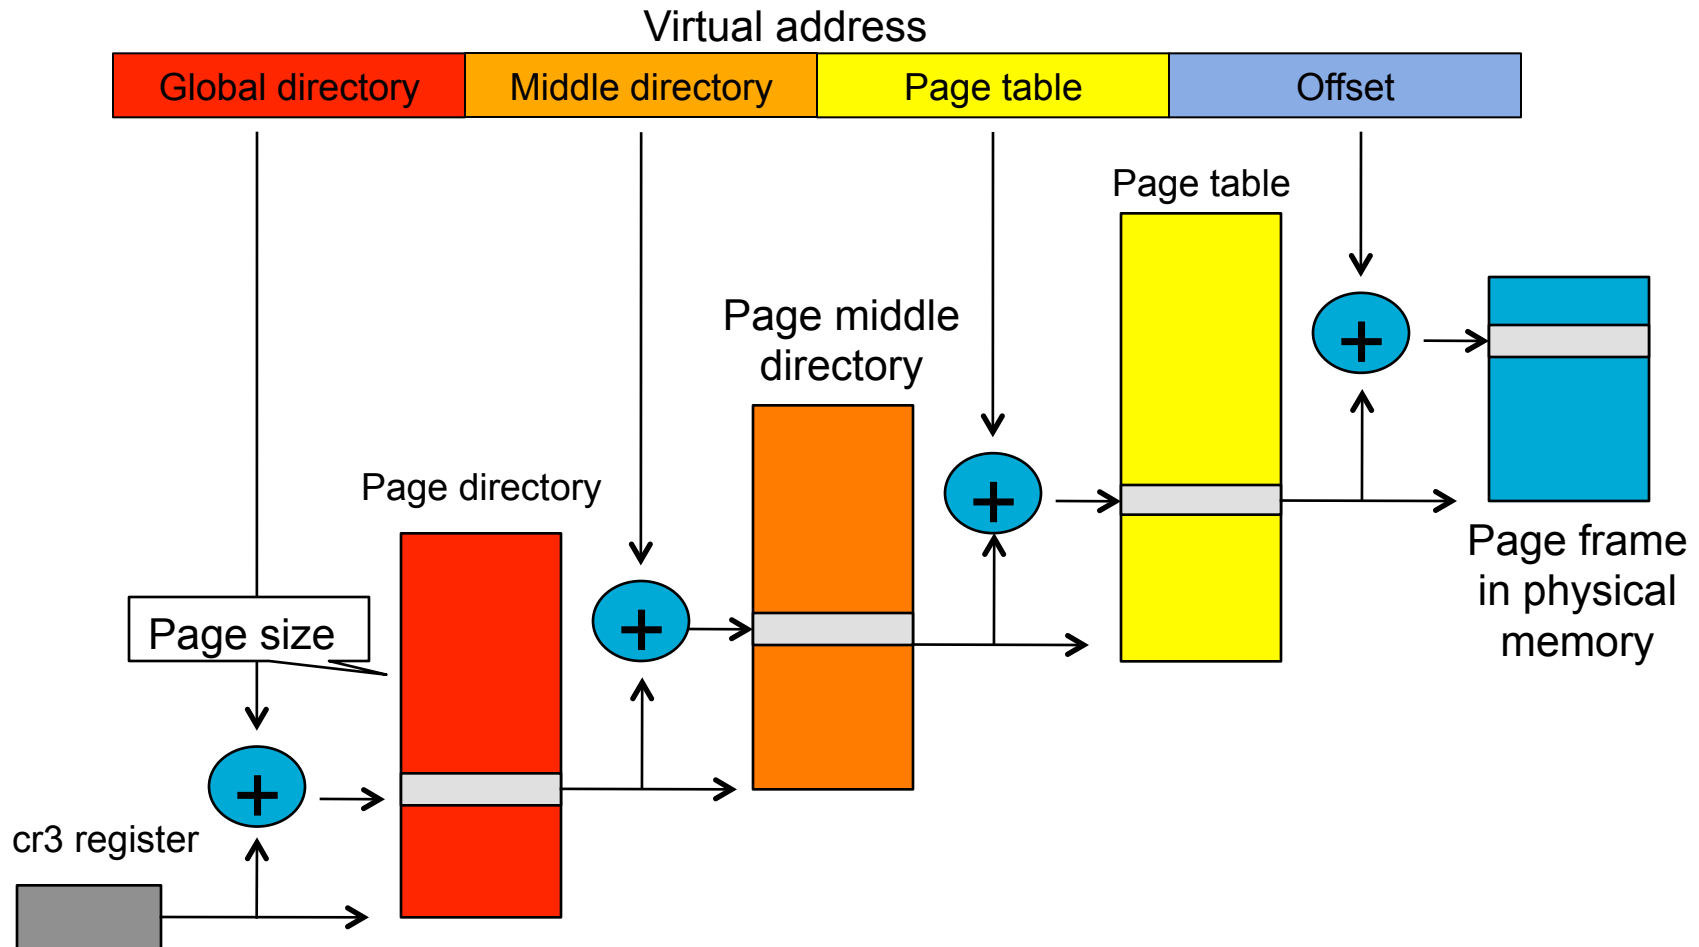
# Hierarchical page table





# Three-level page table in Linux

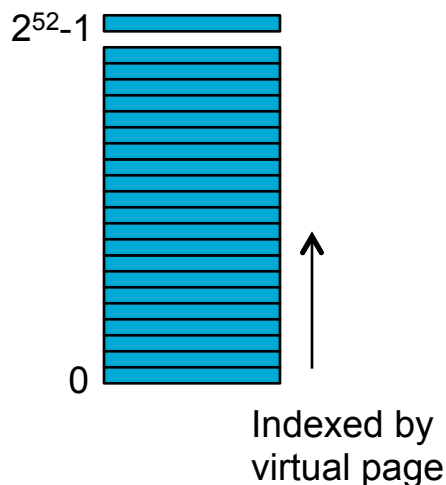
- Designed to accommodate the 64-bit Alpha
  - To adjust for a 32-bit proc. – middle directory of size 1




# Inverted page tables

- Another way to save space – inverted page tables
  - Page tables are indexed by virtual page #, hence their size
  - Inverted page tables – one entry per page frame
    - But to get the page you are still given a VPN
      - Straightforward with a page table, but
    - Linear with inverted page tables – too slow mapping!

Traditional page table  
with an entry per  
each  $2^{52}$  pages




1GB physical  
memory has  $2^{18}$   
4KB page frames




# Inverted and hashed page tables

- Slow inverted page tables ... hash tables may help
  - Different virtual page number might map to identical hash values – use a collision-chain mechanism
  - Problem – poor cache locality

1GB physical  
memory has  $2^{18}$   
4KB page frames



Hash  
table



Indexed by  
hash on  
virtual page

Virtual page | page  
frame

- Also Translation Lookaside Buffers (TLB) ...

*And now a short break ...*

---

# Speeding things up a bit

---

- Simple page table 2x cost of memory lookups
  - First into page table, a second to fetch the data
- Two-level page tables triple the cost!
  - Two lookups into page table and then fetch the data
- And two is not enough ...
- *How can we make this more efficient?*

# Speeding things up a bit

---

- Ideally, make fetching from a virtual address almost as efficient as from a physical address
- Observation – locality of references (lot of references to few pages)
- Solution – hardware cache inside the CPU
  - Translation Lookaside Buffer
  - Cache the virtual-to-physical translations in the hardware
    - A better name would be address-translation cache
  - Traditionally managed by the MMU

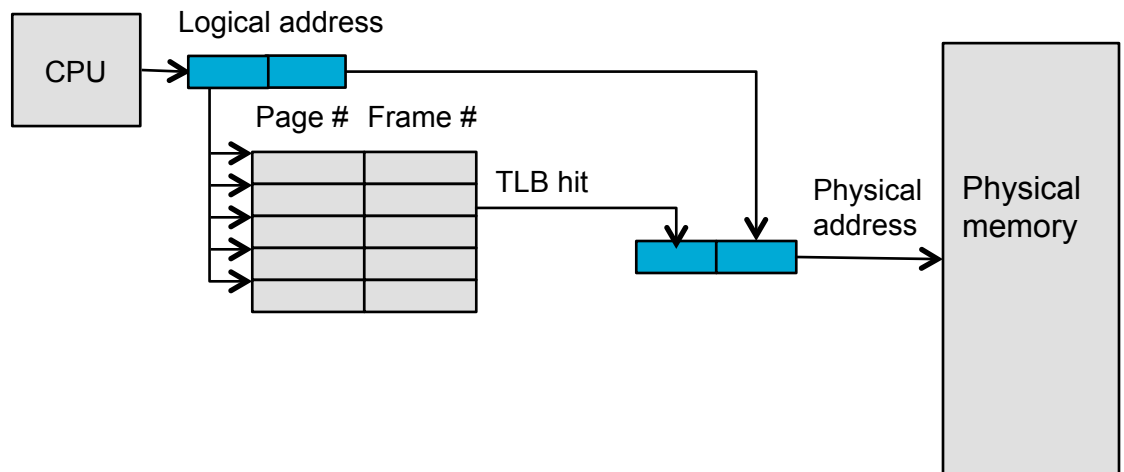
# TLBs

---

- Translates virtual page #s into page frame #s
  - Can be done in single machine cycle
- Implemented in hardware
  - A fully associative cache (parallel search)
  - Cache tags are virtual page numbers
  - Cache values are page frame numbers
    - With this + offset, MMU can calculate physical address
  - A typical TLB entry might look like this  
VPN | PFN | Other bits

# TLBs hit


```
VPN = (VirtAddr * VPN_MASK) >> SHIFT
(Success, TlbEntry) = TLB_Lookup(VPN)
If (Success == True) // TLB Hit
    if (CanAccess(TlbEntry.ProtectBits) == True)
        Offset = VirtAddr & OFFSET_MASK
        PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
        Register = AccessMemory(PhysAddr)
    else
        RaiseException(PROTECTION_FAULT)
```





# TLBs miss

```
else // TLB Miss
    PTEAddr = PTBR + (VPN * sizeof(PTE))
    PTE = AccessMemory(PTEAddr)
    if (PTE.Valid == False)
        RaiseException(SEGMENTATION_FAULT)
    else
        TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
        RetryInstruction()
```



# Managing TLBs

- Address translations mostly handled by TLB
  - >99% of translations, but there are TLB misses
  - If a miss, translation is placed into the TLB
- Who manages the TLB miss?
  - Hardware (memory management unit – MMU)
    - Knows where page tables are in memory
      - OS maintains them, HW access them directly
  - E.g., Intel x86
- Software loaded TLB (OS)
  - TLB miss faults to OS, OS finds page table entry & loads TLB
  - Must be fast
    - CPU ISA has instructions for TLB manipulation
    - OS gets to pick the page table format
  - E.g. MIPS R10k, Sun's SPARC v9

# Managing TLBs

---

- OS must ensure TLB and page tables are consistent
  - When OS changes protection bits in an entry, it needs to invalidate the line if it is in the TLB
- When the TLB misses, and a new process table entry is loaded, a cached entry must be evicted
  - How to choose a victim is called “TLB replacement policy”
  - Implemented in hardware, usually simple (e.g., LRU)


# Managing TLBs

---

- What happens on a process context switch?
  - Remember, each process typically has its own page tables
  - Need to invalidate all the entries in TLB! (flush TLB)
    - A big part of why process context switches are costly
  - *Can you think of a hardware fix to this?*
  - Add an Address Space Identifier field to the TLB

# An example TLB

- From MIPS R4000 – software-managed TLB
  - 32b address space with 4KB pages
    - 20b VPN and 12b offset
  - But TLB has only 19b for VPN!
    - *User addresses will only come from half the address space – the rest is for the kernel*
    - So 19b is enough



# An example TLB

---

● ...

- VPN translates to up to a 24b physical frame number and can thus support systems with up to 64GB of physical memory
- *How is that?!*
- *G* is for pages globally shared (so ASID is ignored)
- *C* is for coherence; *D* is for dirty and *V* is for valid
- Since it is software managed, OS needs instructions to manipulate it
  - TLBP – probes
  - TLBR – reads
  - TLBWI and TLBWR to replaces a specific or a random entry

# Effective access time

- Associative Lookup =  $\epsilon$  time units
- Hit ratio -  $\alpha$  - percentage of times that a page number is found in the associative registers (ratio related to TLB size)

Effective Memory Access Time (EAT)

$$\text{EAT} = \alpha * (\epsilon + \text{memory-access}) + (1 - \alpha) (\epsilon + 2 * \text{memory-access})$$

TLB hit

TLB miss

$$\alpha = 80\%$$

$$\epsilon = 20 \text{ nsec}$$

$$\text{memory-access} = 100 \text{ nsec}$$

$$\text{EAT} = 0.8 * (20 + 100) + 0.2 * (20 + 2 * 100) = 140 \text{ nsec}$$

Why 2?

# Next time

---

- Virtual memory policies
- Some other design issues