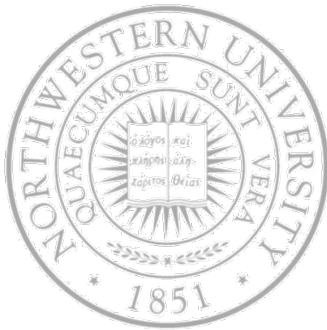


Synchronization I



Today

- Race condition, critical regions and locks

Next time

- Condition variables, semaphores and Monitors

Cooperating processes

- Cooperating processes need to communicate
 - They can affect/be affected by others
- Issues with Inter-Process Communication (IPC)
 1. How to pass information to another process?
 2. How to avoid getting in each other's ways?
 - Two processes trying to get the last seat on a plane
 3. How to ensure proper sequencing when there are dependencies?
 - Process A produces data that B prints – B must wait for A to print
- How about threads?
 - 1. Easy
 - 2 & 3. Pretty much the same

Accessing shared resources

- Many times cooperating threads share memory
- A common example – print spooler
 - A thread wants to print a file, enter file name in a special spooler directory
 - Printer daemon, another thread, periodically checks the directory, prints whatever file is there and removes the name

```
next_slot:= in;                // in = 4
spooler_dir[next_slot] := file_name; // insert "abc"
in++;
```

Spooler
directory

3	zzz
4	abc
5	
6	
7	
8	

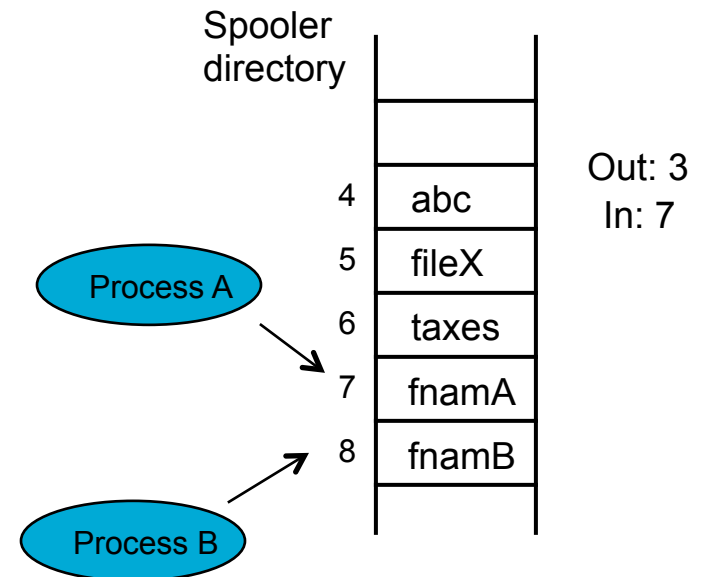
Out: 3
In: 4

Accessing shared resources

- Assumption – preemptive scheduling
- Two threads, A & B, trying to print

Switch

```
A: next_slot_A ← in          % 7
A: spooler_dir[next_slot_A] ← file_name_A
A: in ← next_slot_A + 1      % 8
-----
B: next_slot_B ← in          % 8
B: spooler_dir[next_slot_B] ← file_name_B
B: in ← next_slot_B + 1      % 9
```



Interleaved schedules

- Problem – the execution of the two threads/processes can be interleaved
 - Some times the result of interleaving is OK, other times not!

Switch $A: next_slot_A \leftarrow in \quad \% 7$

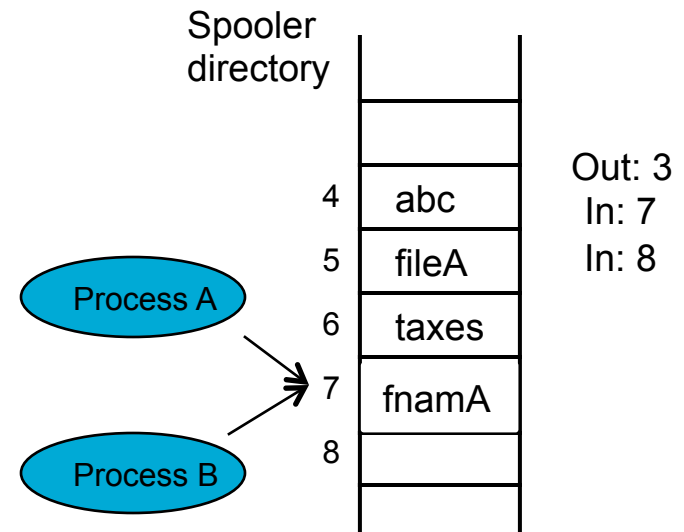
 $B: next_slot_B \leftarrow in \quad \% 7$

Switch $B: spooler_dir[next_slot_B] \leftarrow file_name_B$ -

Switch $B: in \leftarrow next_slot_B + 1 \quad \% 8$

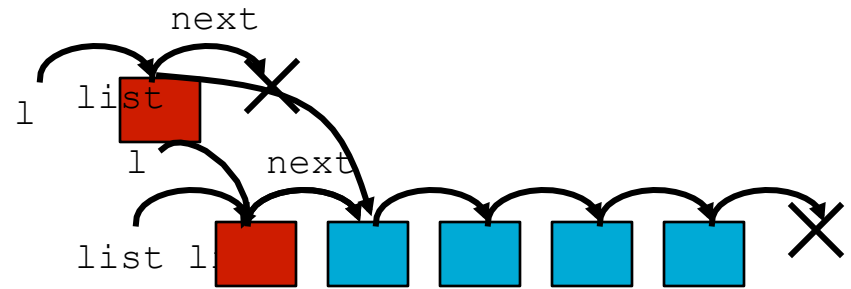
 $A: spooler_dir[next_slot_A] \leftarrow file_name_A$
 $A: in \leftarrow next_slot_A + 1 \quad \% 8$

B's printout?!



Interleaved schedule – another example

```
1 struct list {  
2   int data;  
3   struct list *next;  
4 };  
5  
6 struct list *list = 0;  
7  
8 void  
9 insert(int data)  
10 {  
11   struct list *l;  
12  
13   l = malloc(sizeof *l);  
14   l->data = data;  
15   l->next = list;  
16   list = l;  
17 }
```



Two threads A and B, what would happen if B executes line 15 before A executes 16?

Race conditions and critical regions

- Problem – the threads operating on the data assumes certain conditions (invariants) hold
 - For the linked list – `list` points to the head of the list and each element's `next` point to the next element
 - Insert temporarily violates this, but fixes it before finishing
 - *True for a single thread, not for two concurrent ones*
- *Race condition*
 - Two or more threads/processes access (r/w) shared data
 - Final results depends on order of execution
- Code where race condition is possible – *critical region*

Race conditions and critical regions

- We need mechanisms to prevent race conditions, synchronizing access to shared resources
 - Some tools try to detect them – `helgrind`
- We need a way to ensure the invariant conditions hold when the process is going to manipulate the share item, i.e. ...
- ... to ensure that *if a thread is using a shared item, others will be excluded from doing it*
 - i.e. *only one thread at a time in the critical region (CR)*

Mutual exclusion

Requirements for a solution

1. No two threads simultaneously in CR
 - Mutual exclusion, at most one thread in
2. No assumptions on speeds or numbers of CPUs
3. No thread outside its CR can block another one
 - Ensure progress; a thread outside the CR cannot prevent another one from entering
4. No thread should wait forever to enter its CR
 - Bounded waiting or no starvation
 - Threads waiting to enter a CR should *eventually* be allow to enter

How about taking turns?

- Strict alternation

- `turn` keeps track of whose turn it is to enter the CR

Thread 0

```
while(TRUE) {  
    while(turn != 0);  
    critical_region0();  
    turn = 1;  
    noncritical_region0();  
}
```

Thread 1

```
while(TRUE) {  
    while(turn != 1);  
    critical_region1();  
    turn = 0;  
    noncritical_region1();  
}
```

- Problems?

- What if thread 0 sets `turn` to 1, but it gets around to just before its critical region before process 1 even tries?
 - `Turn` is 1 and both process are in their noncritical region
 - Violates conditions 3

Locks

- Using locks
 - It's a variable so you have to declare it
 - Threads check lock when entering CR, and free it after
 - Lock is either available (free) or acquired
 - Can hold other information such as which thread holds the lock or a queue of lock requests
 - `lock()`: if available, go on; else don't return until you have it
 - `unlock()`: if threads are waiting, they will (eventually) find out (or be told)

```
lock_t mutex;
```

```
void  
insert(int data)  
{  
    struct list *l;  
  
    lock(&mutex);  
    l = malloc(sizeof *l);  
    l->data = data;  
    l->next = list;  
    list = l;  
    unlock(&mutex);  
}
```

Pthreads locks

- In the POSIX library, a lock is called a **mutex**

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;

void
insert(int data)
{
    struct list *l;

    Pthread_mutex_lock(&lock); /* a wrapper that checks for errors */
    l = malloc(sizeof *l);
    ...
    Pthread_mutex_unlock(&lock);
}
```

- Note the call passes a variable to lock/unlock
 - To enable fine-grain locking (rather than a single coarse lock)
- Locks must be initialized before used, either this way or dynamically with

```
pthread_mutex_init(&lock, NULL)
```

Implementing locks

- Here it's a simple implementation of `lock()`
- *Are we done?*

```
1 void
2 lock(lock *lk)
3 {
4     while(lk->locked == 1)
5         ; /* spin-wait */
6     lk->locked = 1; /* now set it */
7 }
```

Context switch
here and there
we go again

- No yet!
 - Correctness problem: Both can concurrently test 4, see it unlocked, and grab it; now both are in the CR
 - Continuously testing a variable for a given value is called *busy waiting*; a lock that uses this is a *spin lock* – spin waiting is wasteful

Implementing locks

- Disabling interrupts

- Simplest solution – threads disables all interrupts when entering the CR and re-enables them at exit

```
void lock() {  
    DisableInterrupts();  
}
```

```
void unlock() {  
    EnableInterrupts();  
}
```

- No interrupts → no clock interrupts → no other process getting in your way

- Obvious problems

- Users in control – grabs the CPU and never comes back
- What about multiprocessors?
- And yes, it's also inefficient

- Use in the kernel – still multi-core means we need something more sophisticated

TSL(test&set) -based solution

- Atomically test & modify the content of a word – TSL
 - The CPU executing the TSL locks the memory bus to stop other CPUs from accessing memory until it is done
 - In SPARC is `ldstwb` (load & store), in x86 is `xchg`

```
int TSL(int *ptr, int new) {  
    int old = *ptr; /* fetch old value at ptr */  
    *ptr = new; /*store new value into ptr */  
    return old;  
}
```

} Done atomically

TSL(test&set) -based solution

- Entering and leaving CR

```
typedef struct __lock_t {
    int flag;
} lock_t;

void init(lock_t *loc) {
    loc->flag = 0;
}

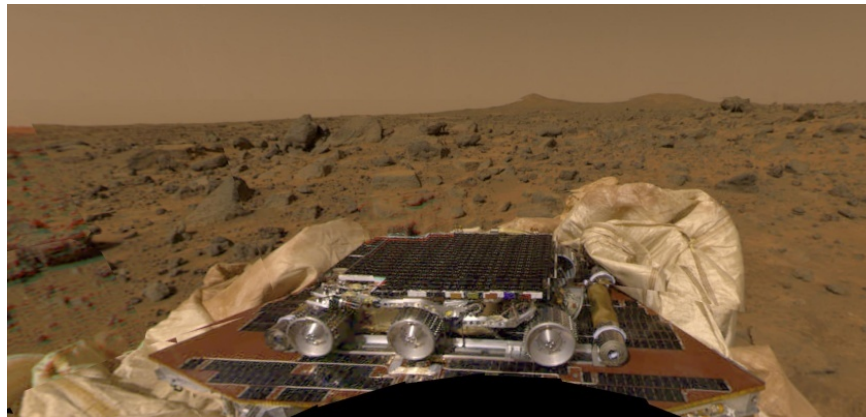
void lock(lock_t *lock) {
    while (TSL(&lock->flag, 1) == 1)
        ; /* spin-wait */
}

void unlock(lock_t *lock) {
    lock->flag = 0;
}
```


Busy waiting and priority inversion

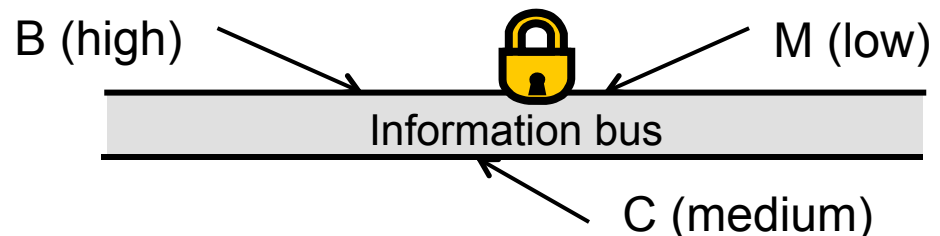
- Problems with TSL-based approach?
 - Waste CPU by busy waiting
 - Can lead to *priority inversion*
 - Two processes, H (high-priority) & L (low-priority)
 - L gets into its CR
 - H is ready to run and starts busy waiting
 - L is never scheduled while H is running ...
 - *So L never leaves its critical region and H loops forever!*

Welcome to Mars!



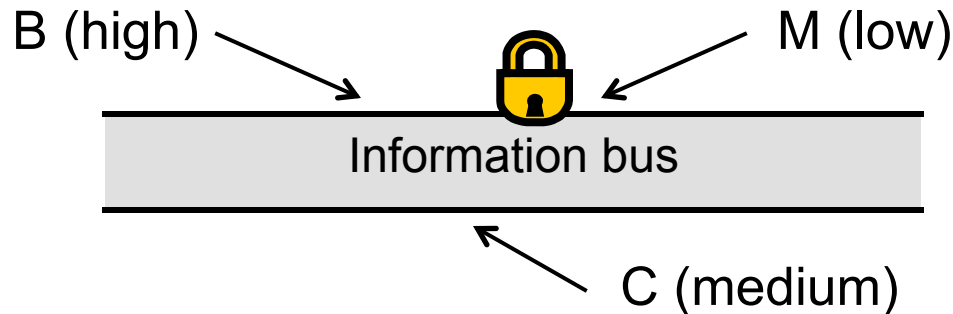
Problems in the Mars Pathfinder*

- Mars Pathfinder
 - Launched Dec. 4, 1996, landed July 4th, 1997
- Periodically the system reset itself, losing data
- Pathfinder software architecture
 - An information bus with access controlled by a lock
 - A bus management (B) high-priority thread
 - A meteorological (M) low-priority, short-running thread
 - If B thread was scheduled while M thread was holding the lock, B busy waited on the lock
 - A communication (C) thread running with medium priority



*As explained by D. Wilner, CTO of Wind River Systems, and narrated by Mike Jones

Problems in the Mars Pathfinder*



- Sometimes,
 - **B (high-priority) was waiting on M (low priority) and**
 - **C (medium priority) was scheduled**
- After a bit of waiting, a watchdog timer would reset the system 😊
- How would you fix it?
 - Priority inheritance – the M thread inherits the priority of the B thread blocked on it
 - Actually supported by VxWork but disabled!

Yield rather than spin

```
void lock(lock_t *lock) {  
    while (TSL(&lock->flag, 1) == 1)  
        ; /* spin-wait */  
}
```

- Too much spinning
 - Imagine two threads; first one gets the lock and is interrupted
 - Second one wants the lock and have to wait ... and wait ...
 - Rather than sit in a tight loop, go to sleep
- An alternative – just yield

```
void lock(lock_t *lock) {  
    while (TSL(&lock->flag, 1) == 1)  
        yield(); /* give up the CPU */  
}
```

- Better than spinning but
 - What about the context switching cost?
 - Is there a chance of starvation?

Sleep rather than spin

- What if the *wrong* thread is waken up?
 - The one not holding the lock – wasted context switch
- Too much left to chance
 - The schedule determines who runs next; if it makes a bad choice – yield immediately or sleep
 - Let's get some control over who gets to acquire the lock next

```
typedef struct __lock_t {  
    int flag;  
    int guard;  
    queue_t *q;  
} lock_t;
```

An explicit queue
to control who
gets the lock

```
void init(lock_t *m) {  
    m->flag = 0;  
    m->guard = 0;  
    queue_init(m->q);  
}
```

Sleep rather than spin

- Two special calls (from Solaris)
 - `park()` / `unpark()` – put calling thread to sleep / wake one up

```
void lock(lock_t *m) {
    while (TSL(&m->guard, 1) == 1)
        ;
    if (m->flag == 0) {
        m->flag = 1;
        m->guard = 0;
    } else {
        queue_add(m->q, gettid());
        m->guard = 0;
        park();
    }
}
```

Note the use of guard as a spin-lock around flag and queue manipulation – so not quite avoiding spinning

Here it's where the thread is when woken up

```
void unlock(lock_t *m) {
    while (TSL(&m->guard, 1) == 1)
        ;
    if (queue_empty(m->q))
        m->flag = 0;
    else
        unpark(queue_remove(m->q));
    m->guard = 0;
}
```

Notice we are not setting flag back to zero when waking up a thread; the one woken up does not have it

Sleep rather than spin

```
void lock(lock_t *m) {
    while (TSL(&m->guard, 1) == 1)
        ;
    if (m->flag == 0) {
        m->flag = 1;
        m->guard = 0;
    } else {
        queue_add(m->q, gettid());
        m->guard = 0;
        park();
    }
}
```

Just curious, what would happen if you park before releasing guard?

Isn't that a race condition? What would happen if the thread about to park is interrupted, the one holding the lock releases it ... the parking one will never wakeup!

- One solution uses a third Solaris system call
 - `setpark()` – I am about to park, so be ware
 - After this, if the thread is interrupted and another calls `unpark` before the `park` is called, `park` returns immediately

```
...
} else {
    queue_add(m->q, gettid());
    setpark();
    m->guard = 0;
    park();
}
```

Lock-based concurrent data structures

- Making data structures *thread-safe*, i.e., usable by threads
 - Two concerns – correctness, obviously
 - And performance
- A non-concurrent counter

```
typedef struct __counter_t {
    int value;
} counter_t;

void init(counter_t *c) {
    c->value = 0;
}

void increment(counter_t *c) {
    c->value++;
}

void decrement(counter_t *c) {
    c->value--;
}
```


Lock-based concurrent data structures

- A concurrent version

```
typedef struct __counter_t {  
    int value;  
    pthread_lock_t lock;  
} counter_t;  
  
void init(counter_t *c) {  
    c->value = 0;  
    pthread_mutex_init(&c->lock, NULL);  
}
```

```
void increment(counter_t *c) {  
    pthread_mutex_lock(&c->lock);  
    c->value++;  
    pthread_mutex_unlock(&c->lock);  
}  
  
void decrement(counter_t *c) {  
    pthread_mutex_lock(&c->lock);  
    c->value--;  
    pthread_mutex_unlock(&c->lock);  
}
```

- Still not very scalable; for a better option with sloppy counters
 - S. Boyd-Wickizer et al., “An analysis of Linux Scalability to Many Cores,” OSDI 2010

Lock-based concurrent data structures

- A first concurrent list (part of, actually)

```
typedef struct __node_t {
    int key;
    struct __node_t *next;
} node_t;

typedef struct __list_t {
    node_t *head;
    pthread_mutex_t lock;
} list_t;
```

```
void List_Init(list_t *L) {
    L->head = NULL;
    pthread_mutex_init(&L->lock, NULL);
}
```

```
int List_Insert(list_t *L, int key) {
    pthread_mutex_lock(&L->lock);
    node_t *new = malloc(sizeof(node_t));
    if (new == NULL) {
        perror("malloc");
        pthread_mutex_unlock(&L->lock);
        return -1; /* failure */
    }
    new->key = key;
    new->next = L->head;
    L->head = new;
    pthread_mutex_unlock(&L->lock);
    return 0; /* success */
}
```

```
int List_Lookup(list_t *L, int key) {
    pthread_mutex_lock(&L->lock);
    node_t *current = L->head;
    while(curr) {
        if (curr->key == key) {
            pthread_mutex_unlock(&L->lock);
            return 0; /* success */
        }
        curr = curr->next;
    }
    pthread_mutex_unlock(&L->lock);
    return -1; /* failure */
}
```

Can we simplify this to avoid releasing the lock on the failure path?

Very coarse locking; what's your critical section?

Lock-based concurrent data structures

- And some improvements

```
int List_Insert(list_t *L, int key) {
    pthread_mutex_lock(&L->lock);
    node_t *new = malloc(sizeof(node_t));
    if (new == NULL) {
        perror("malloc");
        pthread_mutex_unlock(&L->lock);
        return -1; /* failure */
    }
    new->key = key;
    new->next = L->head;
    L->head = new;
    pthread_mutex_unlock(&L->lock);
    return 0; /* success */
}
```

```
int List_Lookup(list_t *L, int key) {
    pthread_mutex_lock(&L->lock);
    node_t *current = L->head;
    while(curr) {
        if (curr->key == key) {
            pthread_mutex_unlock(&L->lock);
            return 0; /* success */
        }
        curr = curr->next;
    }
    pthread_mutex_unlock(&L->lock);
    return -1; /* failure */
}
```

```
int List_Insert(list_t *L, int key) {
    node_t *new = malloc(sizeof(node_t));
    if (new == NULL) {
        perror("malloc");
        return -1; /* failure */
    }
    new->key = key;
    pthread_mutex_lock(&L->lock);
    new->next = L->head;
    L->head = new;
    pthread_mutex_unlock(&L->lock);
    return 0; /* success */
}
```

Just lock the critical section

```
int List_Lookup(list_t *L, int key) {
    int rv = -1;
    pthread_mutex_lock(&L->lock);
    node_t *current = L->head;
    while (curr) {
        if (curr->key == key) {
            rv = 0;
            break; /*
        }
        curr = curr->next;
    }
    pthread_mutex_unlock(&L->lock);
    return rv;
}
```

A single return path, to avoid potential bugs

Coming up ...

- Other mechanisms for synchronization
 - Condition variables
 - Semaphores – slightly higher abstractions
 - Monitors – much better but requiring language support