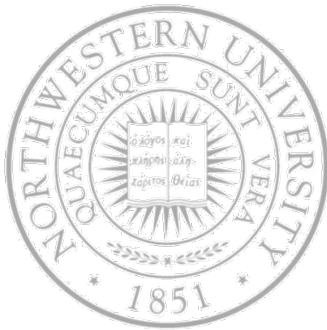


Synchronization II



Today

- Condition Variables
- Semaphores
- Monitors
- and some classical problems

Next time

- Deadlocks

Condition variables

- Many times a thread wants to check whether a condition is true before continuing execution
 - A parent waiting on a child
 - A consumer waiting on something to consume
 - Spinning on a shared variable is inefficient
- Condition variables!
 - An explicit queue where threads can put themselves on when some state is not what they want (*waiting* on a change)
 - Until some other thread changes the state and informs them of it, *signaling* on the condition

```
pthread_cond_t c;  
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);  
pthread_cond_signal(pthread_cond_t *c);
```

Waiting on your child

- Before an example, did you notice?

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);
```

- Assumes mutex is locked before wait is called
 - Wait must release it and put the thread to sleep, *atomically*
 - When the thread wakes up, re-acquires the lock before returning
- All this to prevent race conditions when the thread is trying to put itself to sleep

- Back to parent and child

```
int done = 0;
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t c = PTHREAD_COND_INITIALIZER;
```

```
int main(int argc, char *argv[]) {
    pthread_t p;
    printf("parent: begin\n");
    pthread_create(&p, NULL, child, NULL);
    thr_join();
    printf("parent: done\n");
    return 0;
}
```

Waiting on your child

```
void thr_exit() {  
    pthread_mutex_lock(&m);  
    done = 1;  
    pthread_cond_signal(&c);  
    pthread_mutex_unlock(&m);  
}  
  
void *child(void *arg) {  
    printf("child\n");  
    thr_exit();  
    return 0;  
}
```

```
void thr_join() {  
    pthread_mutex_lock(&m);  
    while (done == 0)  
        pthread_cond_wait(&c, &m);  
    pthread_mutex_unlock(&m);  
}
```

That while doesn't seem strictly necessary, wouldn't an if do ... wait a bit

- Two cases to consider
 - Parent creates the child and continue running
 - Gets the lock, check if done and put itself to sleep
 - Child eventually runs, gets the lock, sets done and signals the parent
 - Parent returns from the wait with the lock held, unlocks it and is done
 - If child runs first, sets done, signals (nobody is waiting) and returns; parent check child is done and returns

Non-working approaches

```
void thr_exit() {  
    pthread_mutex_lock(&m);  
    /* done = 1; */  
    pthread_cond_signal(&c);  
    pthread_mutex_unlock(&m);  
}
```

```
void thr_join() {  
    pthread_mutex_lock(&m);  
    /* while (done == 0) */  
    pthread_cond_wait(&c, &m);  
    pthread_mutex_unlock(&m);  
}
```

```
void thr_exit() {  
    done = 1;  
    pthread_cond_signal(&c);  
}  
  
void thr_join() {  
    if (done == 0)  
        pthread_cond_wait(&c);  
}
```

Do you need done?

- If the child runs immediately, the signal will be lost
- Parent will call wait (there's nothing to check) and go to sleep for ever

Do you need that mutex?

- What would happen if the parent is interrupted after checking 'done' but before going to sleep on wait?
- Child runs, signals nobody (parent is not there yet!) and ..
- When parent continues it goes to sleep, *for ever!*


Producer/consumer problem

- Producer-consumer problem, aka *bounded buffer*
 - Two or more processes & one shared, fixed-size buffer
 - Some put data times into a buffer, others takes them out
 - E.g., Web server with producers taken orders and consumer threads processing them

```
int buff[MAX];
int fill = 0;
int use = 0;
int count = 0;
```

```
void put(int value) {
    buff[fill] = value;
    fill = (fill + 1) % MAX;
    count++;
}
```

```
int get() {
    int tmp = buffer[use];
    use = (use + 1) % MAX;
    count--;
    return tmp;
}
```



Producer/consumer problem

- “Simple solution”
 - If buffer empty, producer goes to sleep to be awoken when the consumer has removed one or more items
 - Similarly for the consumer
 - (a first try)

```
cond_t cond;  
mutex_t mutex;
```

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        pthread_mutex_lock(&mutex);  
        if (count == MAX)  
            pthread_cond_wait(&cond, &mutex);  
        put(i);  
        pthread_cond_signal(&cond);  
        pthread_mutex_unlock(&mutex);  
    }  
}
```

```
void *consumer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        pthread_mutex_lock(&mutex);  
        if (count == 0)  
            pthread_cond_wait(&cond, &mutex);  
        int tmp = get(i);  
        pthread_cond_signal(&cond);  
        pthread_mutex_unlock(&mutex);  
    }  
}
```

A while for an if

```
cond_t cond;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        pthread_mutex_lock(&mutex);
        if (count == MAX)
            pthread_cond_wait(&cond, &mutex);
        put(i);
        pthread_cond_signal(&cond);
        pthread_mutex_unlock(&mutex);
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        pthread_mutex_lock(&mutex);
        if (count == 0)
            pthread_cond_wait(&cond, &mutex);
        int tmp = get(i);
        pthread_cond_signal(&cond);
        pthread_mutex_unlock(&mutex);
    }
}
```

2 consumers/1 producer

- Consumer 1 tries to get item but finds buffer empty, signals producer
- Producer puts an item and signals this, moving C1 to ready queue
- Consumer 2 comes along and gets the one item
- Now C1 runs; just before returning from the wait it re-acquires the lock, returns and calls get to find an empty buffer!!

With condition variables, always use while loops

One condition variable is not enough

```
cond_t cond;  
mutex_t mutex;
```

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        pthread_mutex_lock(&mutex);  
        while (count == MAX)  
            pthread_cond_wait(&cond, &mutex);  
        put(i);  
        pthread_cond_signal(&cond);  
        pthread_mutex_unlock(&mutex);  
    }  
}
```

```
void *consumer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        pthread_mutex_lock(&mutex);  
        while (count == 0)  
            pthread_cond_wait(&cond, &mutex);  
        int tmp = get(i);  
        pthread_cond_signal(&cond);  
        pthread_mutex_unlock(&mutex);  
    }  
}
```

2 consumers/1 producer
and lets MAX = 1

- Both consumers try to get the item, find buffer empty and go to sleep
- Producer puts item, wakes up a consumer (1) and goes to sleep
- Consumer comes along and gets the one item and signals ...
- *but who!?* Both producer and Consumer 2 are sleeping

Finally a solution

```
cond_t empty, fill;  
mutex_t mutex;
```

Simple solution – two condition variables

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        pthread_mutex_lock(&mutex);  
        while (count == MAX)  
            pthread_cond_wait(&empty, &mutex);  
        put(i);  
        pthread_cond_signal(&fill);  
        pthread_mutex_unlock(&mutex);  
    }  
}
```

Producer waits on “empty”
and signals “fill”

```
void *consumer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        pthread_mutex_lock(&mutex);  
        while (count == 0)  
            pthread_cond_wait(&fill, &mutex);  
        int tmp = get(i);  
        pthread_cond_signal(&empty);  
        pthread_mutex_unlock(&mutex);  
    }  
}
```

Consumers do the opposite – wait
on “fill” and signal “empty”

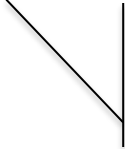
Semaphores

- A synchronization primitive
- Higher level of abstraction than locks, also replacing condition variables
- Invented by Dijkstra in '68 as part of THE operating system
- Atomically manipulated by two operations

```
sem_wait(sem_t *sem) / P / down(sem)  
sem_post(sem_t *sem) / V / up(sem)
```

- The initial value determine its behavior, so it must be first initialized

```
sem_init(sem_t *sem, int pshared, unsigned int value);
```



Ignored this for now, but basically shared by all threads of a process (0) or by processes through shared memory (!=0)

Blocking in semaphores

- Each semaphore has an associated queue of processes/threads
 - sem_wait / P
 - Decrement the value of the semaphore by 1
 - If sem was “unavailable” (negative), wait on the queue
 - P – not really for *proberen* or *passeer* but for a made-up word *prolaag* – “try to reduce”

```
int sem_wait(sem_t *s){  
    s.value--;  
    wait in a queue of s until (s.value > 0);  
}
```

Atomic action

Semaphores

• ...

- sem_post / V
 - Increment the value of the semaphore by one
 - If thread(s) are waiting on the queue, unblock one
- V – *verhogen* – increase in Dutch

```
int sem_post(sem_t *s) {  
    s.value++;  
    if there are 1+ threads waiting  
        wake one thread up;  
}
```

Atomic action

Binary semaphores - locks

```
sem_t m;  
sem_init(&m, 0, 1);
```

Why 1? Look at the definition of wait and post

```
sem_wait(&m);  
/* critical section */  
sem_post(&m);
```

```
int sem_wait(sem_t *s){  
    s.value--;  
    wait in a queue of s until (s.value > 0);  
}
```

```
int sem_post(sem_t *s) {  
    s.value++;  
    if there are 1+ threads waiting  
        wake one thread up;  
}
```

So, if $m = 1$ the first thread will go in and decrement its value, the following thread will wait ... until the thread inside increments it inside post

Semaphores as condition variables

- Waiting on a condition, as when parent waits for child to terminate

```
sem_t s;
```

```
void *child(void *arg) {  
    printf("child\n");  
    sem_post(&s);  
    return NULL;  
}
```

```
int main(int argc, char*argv[])  
    sem_init(&s, 0, 0);  
    printf("parent: begin\n");  
    pthread_t c;  
    pthread_create(&c, NULL, child, NULL);  
    sem_wait(&s);  
    printf("parent: end\n");  
    return 0;  
}
```

Why 0?

So, if $m = 0$ and parent runs, will wait until the child runs and sets value to 1; If child runs first, the value will be 1 and the parent will go on without waiting

Semaphores – Producer/consumer v1

```
sem_t empty;  
sem_t full;  
sem_t mutex;
```

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        put(i);  
        sem_post(&full);  
    }  
}
```

Yeap, those are CSs

```
void *consumer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&full);  
        int tmp = get();  
        sem_post(&empty);  
    }  
}
```

```
int main ...
```

```
...
```

```
sem_init(&empty, 0, MAX); /* MAX buffers are empty ... */
```

```
sem_init(&full, 0, 0);    /* and 0 are full */
```

```
...
```

```
void put(int value) {  
    buff[fill] = value;  
    fill = (fill + 1) % MAX;  
}
```

```
int get() {  
    int tmp = buffer[use];  
    use = (use + 1) & MAX;  
    return tmp;  
}
```


Semaphores – Producer/consumer

```
sem_t empty;  
sem_t full;  
sem_t mutex;
```

```
void *producer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&empty);  
        sem_wait(&mutex);  
        put(i);  
        sem_post(&mutex);  
        sem_post(&full);  
    }  
}
```

Protect the critical section

```
void *consumer(void *arg) {  
    int i;  
    for (i = 0; i < loops; i++) {  
        sem_wait(&full);  
        sem_wait(&mutex);  
        int tmp = get();  
        sem_post(&mutex);  
        sem_post(&empty);  
    }  
}
```

Protect the critical section

```
int main ...  
...  
sem_init(&empty, 0, MAX); /* MAX buffers are empty ... */  
sem_init(&full, 0, 0);    /* and 0 are full */  
sem_init(&mutex, 0, 1);   /* set to 1, it's a lock */  
...
```

Readers-writers problem

- The need for a more flexible type of lock, imagine a database or a simple linked list
 - Not problem with multiple readers allowed at once
 - Only one writer allowed at a time
 - If writers is in, nobody else is

```
typedef struct _rwlock_t {
    sem_t lock;
    sem_t writelock;
    int readers;
} rwlock_t;

void rwlock_acquire_writelock(rwlock_t *rw) {
    sem_wait(&rw->writelock);
}

void rwlock_release_writelock(rwlock_t *rw) {
    sem_post(&rw->writelock);
}
```

Simple, only a single
writer allowed

```
void rwlock_acquire_readlock(rwlock_t *rw) {
    sem_wait(&rw->lock);
    rw->readers++;
    if (rw->readers == 1)
        sem_wait(&rw->writelock);
    sem_post(&rw->lock);
}

void rwlock_release_readlock(rwlock_t *rw) {
    sem_wait(&rw->lock);
    rw->readers--;
    if (rw->readers == 0)
        sem_post(&rw->writelock);
    sem_post(&rw->lock);
}
```

First reader blocks
the writer from entering

Last reader lets
the writer in

Dining philosophers

- Another one by Dijkstra
- Philosophers eat/think
 - To eat, a philosopher needs 2 chopsticks
 - Picks one at a time
- *How to prevent deadlock and starvation*

```
#define N 5
```

```
void philosopher(int i)  
{
```

```
    while (TRUE) {
```

```
        think();  
        take_chopstick(i);  
        take_chopstick((i+1)%N);  
        eat();  
        put_chopstick(i);  
        put_chopstick((i+1)%N);
```


```
    }
```

```
}
```

Now: Everybody takes
the left chopstick!

Nonsolution

Why not just
protect all this
with a mutex?



Dining philosophers example

state[] – too keep track of philosopher's
(eating, thinking, hungry)
s[] – array of semaphores, one per philosopher

```
void philosopher(int i) {  
    while(TRUE) {  
        think();  
        take_chopstick(i);  
        eat();  
        put_chopstick(i);  
    }  
}
```

```
void take_chopstick(int i) {  
    sem_wait(&mutex);  
    state[i] = HUNGRY;  
    test(i);  
    sem_post(&mutex);  
    sem_wait(&s[i]);  
}
```

state

```
void put_chopstick(int i) {  
    sem_wait(&mutex);  
    state[i] = THINKING;  
    test(LEFT);  
    test(RIGHT);  
    sem_post(&mutex);  
}
```

```
void test(int i) {  
    if ((state[i] == hungry &&  
        state[LEFT] != eating &&  
        state[RIGHT] != eating) {  
        state[i] = EATING;  
        sem_post(&s[i]);  
    }  
}
```

Watch for deadlocks

- Semaphores solves most synchronization problems
 - But no control over their use, no guarantee of proper usage

```
void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&full);
        sem_wait(&mutex);
        int tmp = get();
        sem_post(&mutex);
        sem_post(&empty);
    }
}
```

```
void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&mutex);
        sem_wait(&full);
        int tmp = get();
        sem_post(&empty);
        sem_post(&mutex);
    }
}
```

Minor change?

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&empty);
        sem_wait(&mutex);
        put(i);
        sem_post(&mutex);
        sem_post(&full);
    }
}
```

```
void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        sem_wait(&mutex);
        sem_wait(&empty);
        put(i);
        sem_post(&full);
        sem_post(&mutex);
    }
}
```

- Deadlock! Consumer holds the mutex and goes to sleep, to wait for the producer to put something
- Producer can't put anything because consumer holds the lock!

Issues with semaphores

- Solves most synchronization problems, but:
 - We have seen, no control over their use, no guarantee of proper usage (our deadlock example)
 - Also ...
 - Semaphores are essentially shared global variables
 - Can be accessed from anywhere (bad software engineering)
 - No connection bet/ the semaphore & the data controlled by it
 - Used for both critical sections & for coordination (scheduling)

Monitors

- Monitors - higher level synchronization primitive
 - A programming language construct
 - Collection of procedures, variables and data structures
 - Monitor's internal data structures are private
- Monitors and mutual exclusion
 - Only one process active at a time - *how?*
 - Synchronization code is added by the compiler (or the programmer using locks)
- To enforce sequences of events – Condition variables
 - Only accessed from within the monitor
 - Three operations – `wait`, `signal` & `broadcast`

Monitors

- `wait`
 - Atomically releases the lock
 - Suspends execution of the calling thread, place it in the waiting queue
 - When the calling thread is re-enabled, it requires the lock before returning from the wait
- A thread that waits “steps outside” the monitor (to the associated wait queue)
- A condition variable is memoryless, it has no internal state (the shared object defines its own); so, `wait` is not a counter – signal may get lost

Monitors

- `Signal` – Takes one waiting thread off the condition variable's waiting queue and marks it as eligible to run
- `Broadcast` – Like signal but for all threads waiting
- What happen after the signal?
 - Hoare – process awakened run, the other one is suspended
 - Brinch Hansen – process signaling must exit the monitor
 - *Mesa – process signaling continues to run*
- As a programmer – always check the condition after being woken! i.e., call within a loop

```
while (predicateOnStateVar(...)) {  
    wait(&lock);  
}
```

Monitors and semaphores

- Clear similarities between the two – you can use one to implement the other
- A semaphore as a monitor

```
monitor class Semaphore {
    int s;
    Semaphore(int value) {
        s = value;
    }
    void wait() {
        while (s <= 0)
            wait();
        s--;
    }
    void post() {
        s++;
        signal();
    }
};
```

Using it as a binary semaphore

```
Semaphore s(1);

s.wait();
/* Critical section */
s.post();
```

Coming up



- Deadlocks
How deadlock arise and what you can do about them