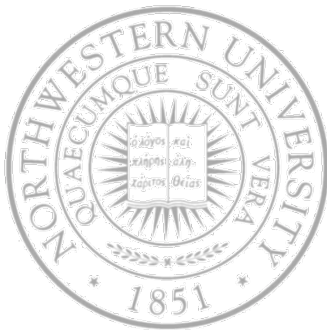


Deadlocks



Today

- Resources & deadlocks
- Dealing with deadlocks
- Other issues

Next Time

- I/O and file systems


"That's some catch, that Catch-22"

Thread A:

```
lock(L1);  
lock(L2);  
...
```

Thread B:

```
lock(L2);  
lock(L1);  
...
```



- *A set of threads is deadlocked if each thread in the set is waiting for an event that only another thread in the set can cause*
- None of the threads can ...
 - run
 - release resources
 - be awakened

Introduction to deadlocks

- Assumptions
 - Threads or single-threaded processes
 - There are no interrupts possible to wake up a blocked thread



- Another “cute” example

“When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up until the other has gone.” An actual law passed by the Kansas legislature ...

Conditions for deadlock

1. Mutual exclusion - Each resource assigned to 1 thread or available
2. Hold and wait - A thread holding resources can request others resources
3. No preemption - Previously granted resources cannot forcibly be taken away
4. Circular wait – A circular chain of 2+ threads, each waiting for resource held by the next thread

All conditions must hold for a deadlock to occur.


Each of the 1-3 conditions is associated with a policy the system can or not have; break one condition → no deadlock

System model


- System – a collection of resources to be shared
- Resources, in types with multiple instances each (printers, files, memory,...)
- Resources can be
 - Preemptable - can be taken away w/o ill effects (e.g. memory)
 - Nonpreemptable - process will fail if resource were taken away (e.g. CD recorder)
- A thread must request a resource before using it & release it once done (`open/close`, `malloc/free`, ...)
 - Sequence of events to use a resource: request/use/release

Deadlock modeling

- Modeled with directed graphs
 - Process B is requesting/waiting for resource S
 - Resource R assigned to process A
 - Process C & D in deadlock over resources T & U




- Generalizing to multiple resource instances per class




Basic facts

- If graph contains no cycles \Rightarrow no deadlock.
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock.
 - if several instances per resource type, maybe a deadlock.



No deadlock here



Deadlock modeling

Clearly, the ordering of operations plays a role

Requests and releases
of each process

A	B	C
Request R	Request S	Request T
Request S	Request T	Request R
Release R	Release S	Release T
Release S	Release T	Release R

... and one particular
ordering


1. A requests R
2. B requests S
3. C requests T
4. A requests S
5. B requests T
6. C requests R

....
deadlock

But with an
alternative

1. A requests R
2. C requests T
3. A requests S
4. C requests R
5. A releases R
6. A releases S

....
no deadlock



Dealing with deadlocks

Possible strategies

- Ignore the problem altogether – ostrich “algorithm”
- Detection and recovery – do not stop it; let it happen, detect it and recover from it
- Dynamic avoidance – careful resource allocation
- Prevention – negating one of the four necessary conditions


The ostrich algorithm

- Pretend there is no problem
- Reasonable if
 - deadlocks occur very rarely
 - cost of prevention is high
- UNIX's & Windows' approach
- A clear trade off between
 - Convenience
 - Correctness
- Not quite an option for your code



Deadlock detection – single instance

- Detect a cycle in the directed graph
 - Simplest case
- *How, when & what*



A linked
list of nodes

1. $L \leftarrow \text{empty}$
all arcs set as unmarked
2. For each node N
/* depth-first search */
 - 2.1. Add N to L & check
if N in L twice there's a
deadlock; exit
 - 2.2. Pick one arc at random,
mark it & follow it to next
current node
3. At end, if no arc no deadlock

Arcs:

$A \rightarrow S, A \leftarrow R, B \rightarrow T, C \rightarrow S$
 $D \rightarrow S, D \leftarrow T, E \rightarrow V, E \leftarrow T$
 $F \rightarrow S, F \leftarrow W, G \rightarrow V, G \leftarrow V$

$L: [R], L: [R, A], L: [R, A, S]$

$L: [B], L: [B, T], L: [B, T, E], \dots$

Detection - multiple instances

n processes, m classes of resources

E – vector of existing resources

A – vector of available resources

C – matrix of currently allocated resources

R – request matrix

C_{ij} – P_i holds C_{ij} instances of resource class j

R_{ij} – P_i wants C_{ij} instances of resource class j

Invariant – $\sum_i C_{ij} + A_j = E_j$
(Currently allocated + available = existing)
i.e. all resources are either allocated or available

Algorithm:

Idea: See if there's any process that can be run to completion with available resources, mark it and free its resources ...

All processes unmarked

1. Look for unmarked process P_i for which $R_i \leq A$
 2. If found, add C_i to A , mark the process and go to 1
 3. If not, exit
- All unmarked processes, if any, are deadlock

Detection

(existing)

$E = (4\ 2\ 3\ 1)$

(available)

$A = (2\ 1\ 0\ 0)$

What process 1 needs

$C =$	<table border="1"><tr><td>0</td><td>0</td><td>1</td><td>0</td></tr></table>	0	0	1	0	$R =$	<table border="1"><tr><td>2</td><td>0</td><td>0</td><td>1</td></tr></table>	2	0	0	1
0	0	1	0								
2	0	0	1								
	2	0	0	1		1	0	1	0		
	0	1	2	0		2	1	0	0		

What process 1 has

Three processes and 4 resource types

After running process 3

$A = (2\ 2\ 2\ 0)$

Now you can run process 2

$A = (4\ 2\ 2\ 1)$

...

Algorithm:

All processes unmarked


1. Look for unmarked process P_i for which $R_i \leq A$
 2. If found, add C_i to A , mark the process and go to 1
 3. If not, exit
- All unmarked processes, if any, are deadlock

When to check & what to do

- *When to try*
 - Every time a resource is requested
 - Every fixed period of times or when CPU utilization drops
- *What to do then - recovery*
 - Through preemption
 - depends on nature of the resource
 - Through rollback
 - Need to checkpoint processes periodically
 - By killing a process
 - Crudest but simplest way to break a deadlock
 - Kill one in or not in the deadlock cycle

Deadlock avoidance

- Dynamically make sure not to get into a deadlock
- Two process resource trajectories
- Every point in the graph, a joint state of the processes



Safe and unsafe states

- Safe if
 - There is no deadlock
 - There is a scheduling order by which all processes can finish
- Un-safe is not deadlock – just no guarantee

Example with one resource (10 instances of it)

Free: 3	Has Needs			Has Needs			Has Needs			Has Needs									
Safe	A	3	9	A	3	9	A	3	9	A	3	9							
	B	2	4	B	4	4	B	0	-	B	0	-							
	C	2	7	C	2	7	C	2	7	C	7	7							
				Free: 1				Free: 5				Free: 0				Free: 7			

Free: 3

Unsafe

	Has	Needs
A	3	9
B	2	4
C	2	7

	Has	Needs
A	4	9
B	2	4
C	2	7

Free: 2

	Has	Needs
A	4	9
B	4	4
C	2	7

Free: 0

	Has	Needs
A	4	9
B	0	-
C	2	7

Free: 4

A requests and is granted another instance

In retrospect, A's request should not have been granted

Banker's algorithm (*Dijkstra, again*)

- Considers
 - Each request as it occurs
 - Sees if granting it leads to a safe state i.e. there are enough resources to satisfy one customer
- With multiple resources
 1. Look for a row $R_i \leq A$, if none the system will eventually deadlock
 2. If found, mark P_i and add C_i to A
 3. Repeat until processes are terminated or a deadlock occurs
- Very cute, but mostly useless
 - Most processes don't know in advance what they need
 - The lists of processes and resources are not static
 - Processes may depend on each other

Deadlock prevention

- Avoidance is pretty hard or impossible
- Can we break one of the condition?
 - Mutual exclusion
 - Hold & wait
 - No preemption
 - Not a viable option
 - How can you preempt a printer?
 - Circular wait

Attacking mutual exclusion

- Some devices can be spooled (printer)
 - Only the printer daemon uses printer resource
 - Thus deadlock for printer eliminated
- But not all devices can be spooled – process table?
- Principle:
 - Assigning resource only when absolutely necessary
 - Reduce number of processes that may claim the resource


Attacking hold & wait

- Processes request all resources at start (~~hold & wait~~)
 - Process never has to wait for what it needs
- But
 - May not know required resources at start
 - It ties up resources others could be using
- Variation (~~hold~~ & wait)
 - Process must release all resources to request a new one

Attacking circular wait

- Impose total order on resources
- Processes request resources in order
- If all processes follow order, no circular wait occurs

Deadlock if $i \rightarrow A \rightarrow j$ & $j \rightarrow B \rightarrow i$
If $i < j$ then $A \rightarrow j \dots$



- Process cannot request resource lower than what it's holding
- Advantage - Simple
- Disadvantage - Arbitrary ordering

Related issues

- Two-phase locking – gather all locks, work & free all
 - If you cannot get all, drop all you have and start again
- Non-resource deadlocks
 - Each is waiting for the other to do some task
 - E.g. communication deadlocks:
 - A sends a request and blocks until B replies, message gets lost!
 - Timeout!
- Livelocks – try, sleep and try again
 - There's some action, just not progress
- Starvation
 - SJF to allocate resources – consider allocation of a printer
 - May cause long job to be postponed indefinitely
 - even though not blocked
 - FIFO?

Next time

- I/O devices, file abstraction and file systems ...