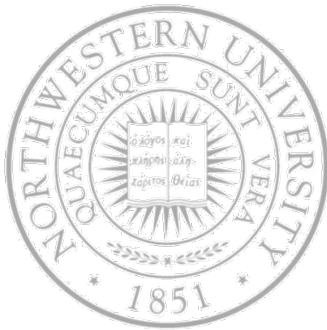


Research in Operating Systems *Sparrow*



Sparrow: Distributed, Low Latency Scheduling

K. Ousterhout, P. Wendell, M. Zaharia and I. Stoica. *In Proc. of SOSP 2013*

(Some) useful concepts picked up from OS

- Scheduler, preemption, shortest-job first, RPC, gang-scheduling, fair-share scheduling, ...

(In no particular order)

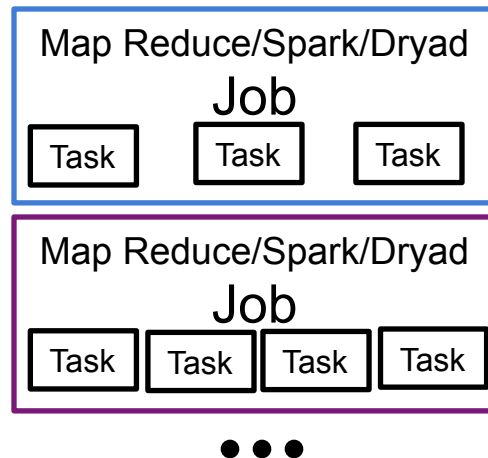
Jobs and scheduling for data analytics

- Large data analytics clusters
- Running ever shorter and higher-fanout jobs
- What for? Finance, language translation, highly personalize search, ...




Jobs and scheduling for data analytics

- Jobs composed of short tasks
- Produced from frameworks that stripe work across 1^3 machines (e.g., Dremel, Spark, ...)
- Targeting task running in $\sim 100\text{ms}$



Job latencies decreasing rapidly



Scheduling challenges

- Millisecond latency
- Quality placement
- Fault tolerant
- High throughput
- *Fixing centralized schedulers is not an option*

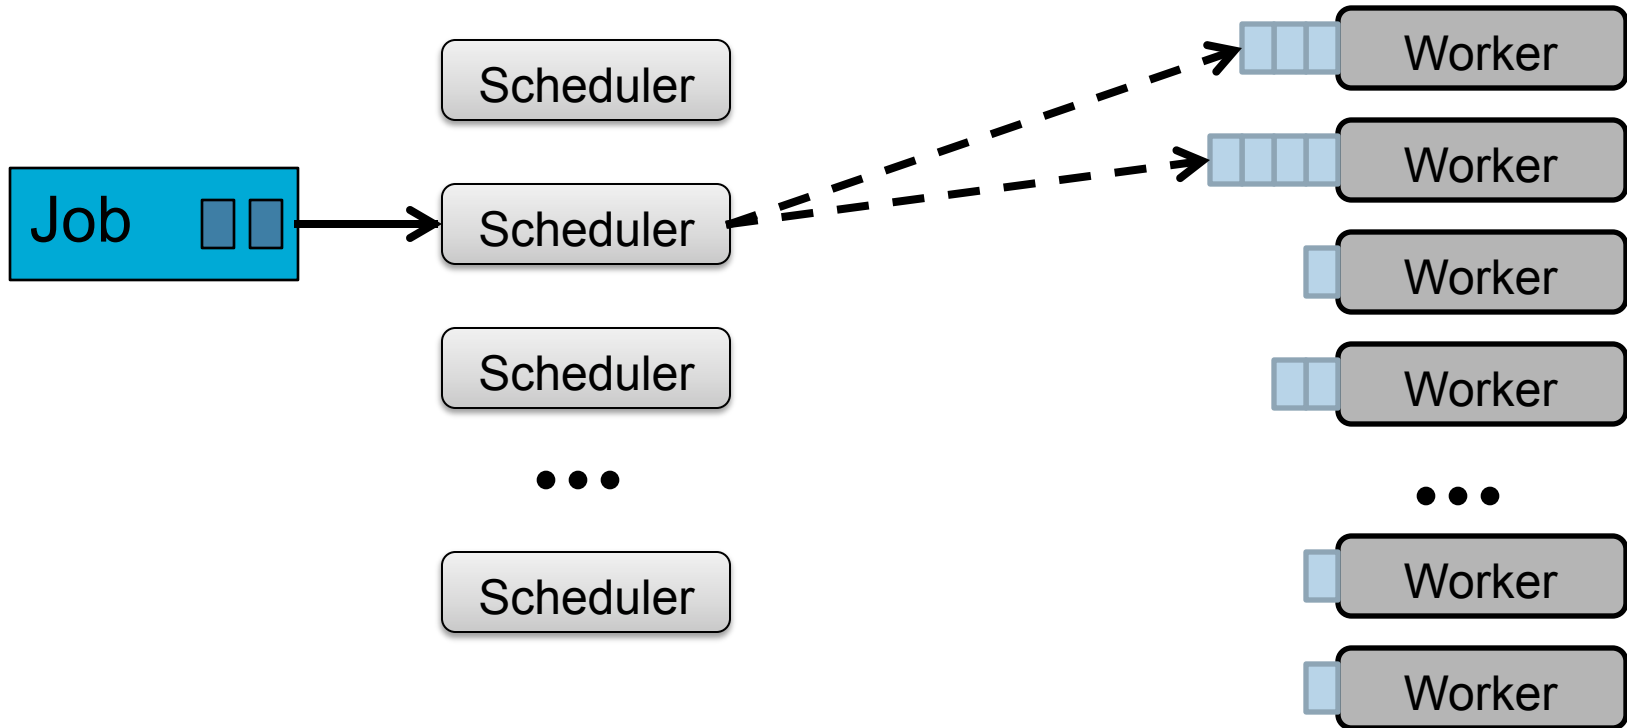
Sparrow schedules tasks in clusters

- using a decentralized, randomized approach
- support constraints and fair sharing and
- provides response times within 12% of ideal
- Adapts power-of-two-choices to parallel task scheduling, introducing three techniques
 - Batch sampling
 - Late binding
 - Policies and constraints

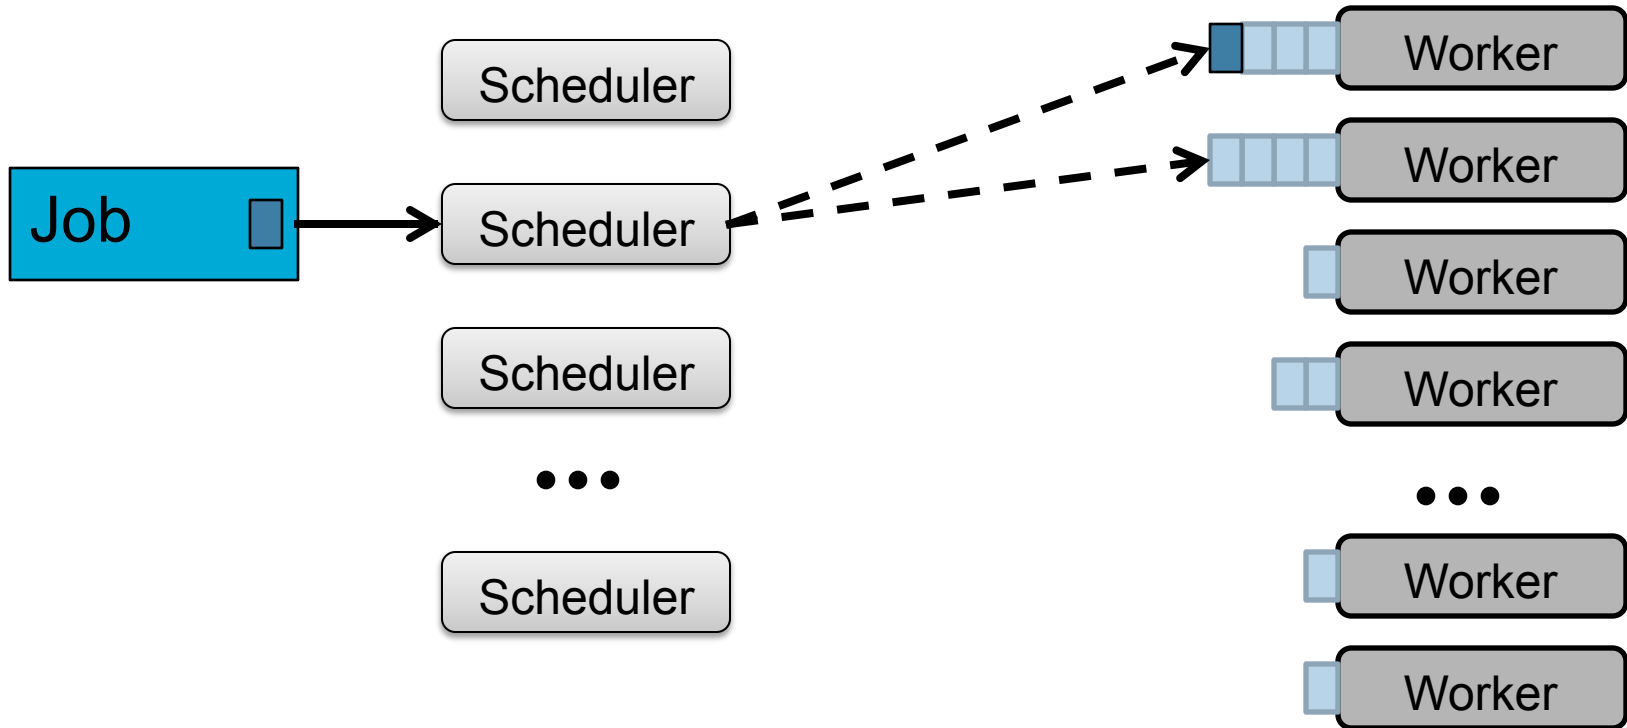
Per-task sampling

- A direct application of power of two choices
- For each task in a job, the scheduler
 - Randomly selects two workers
 - Probes each worker (a lightweight RPC) for queue length
 - Places task in shortest queue

Per-task Sampling



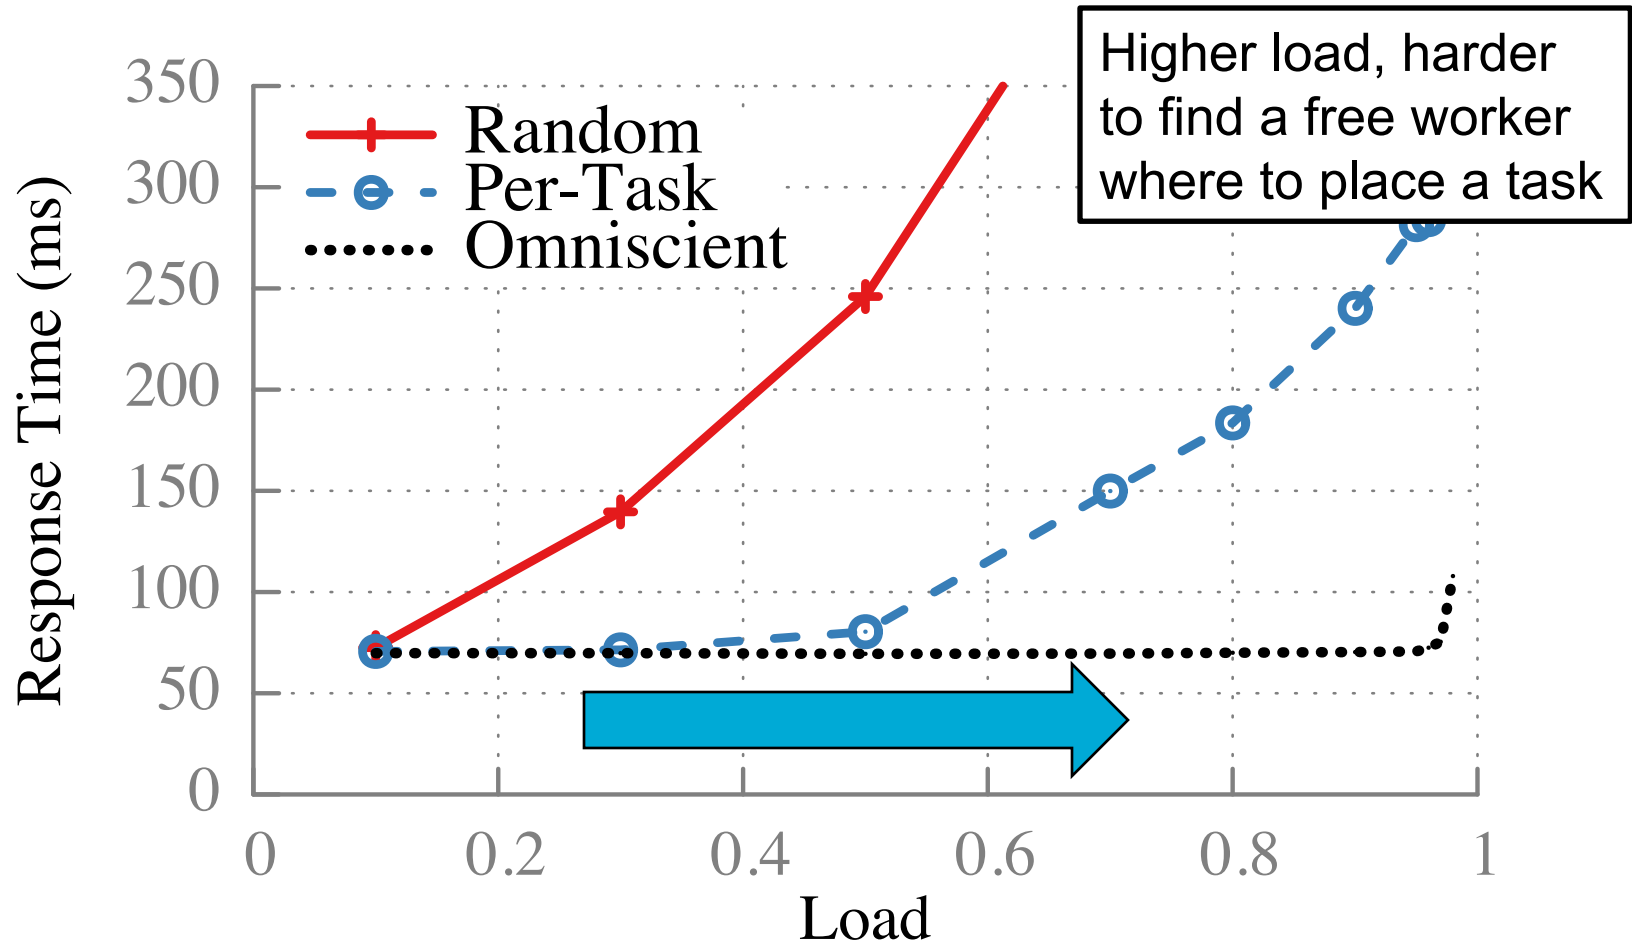
Per-task sampling



Per-task sampling

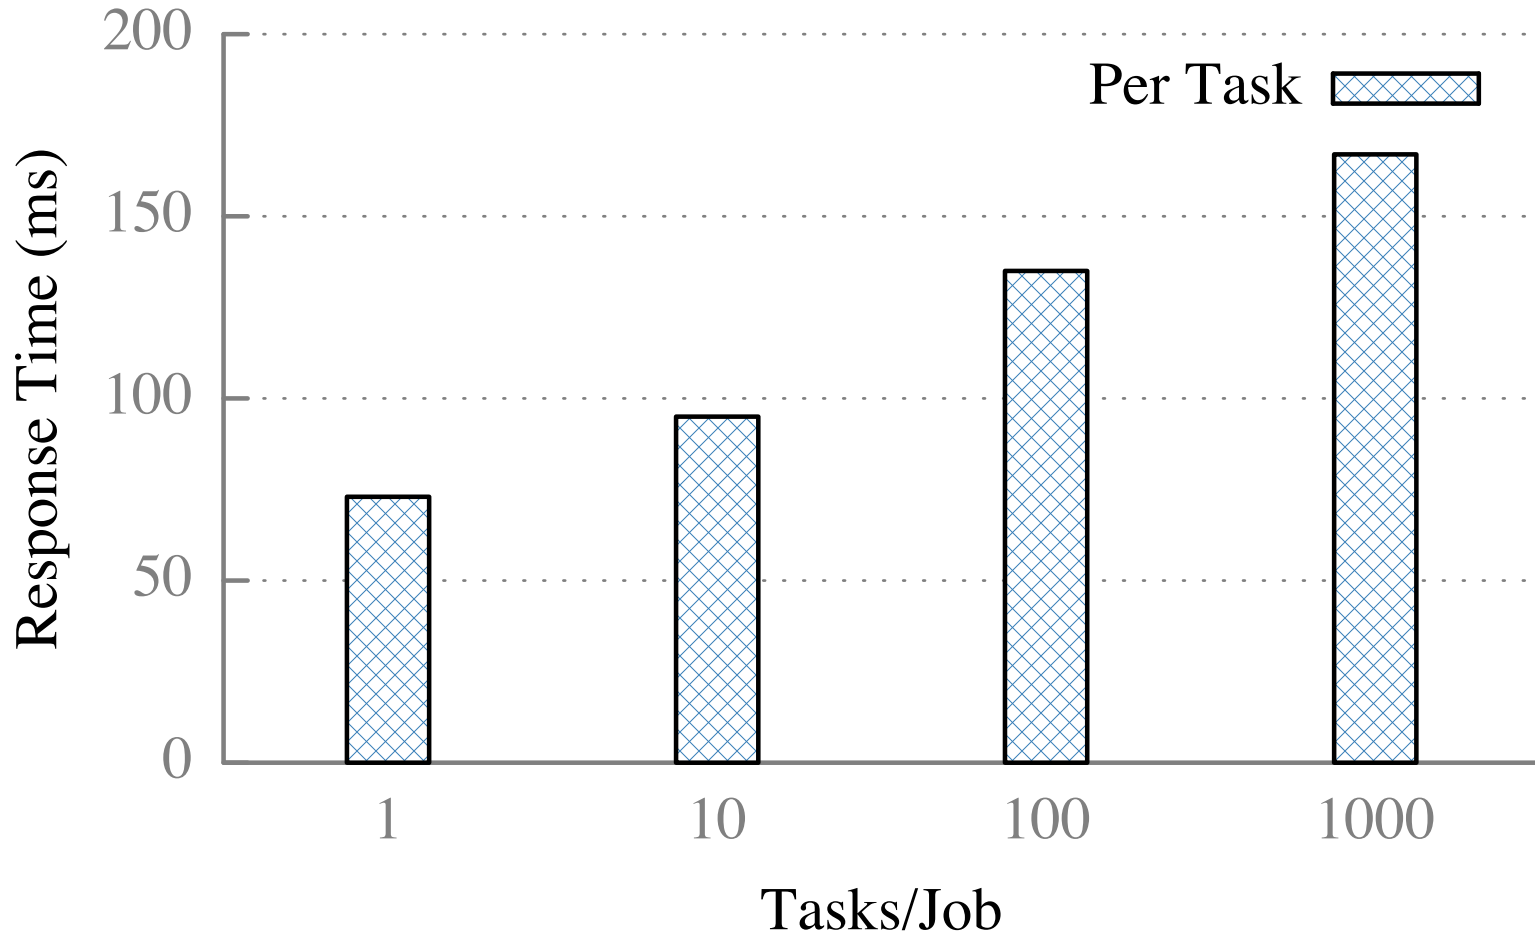
- A direct application of power of two choices
- For each task in a job, the scheduler
 - Randomly selects two workers
 - Probes each worker (a lightweight RPC) for queue length
 - Places task in shortest queue
- For comparison
 - Random
 - Omniscient – greedy, based on complete information

Better than random, $>2\times$ worst than opt



100-task jobs in 10,000-node cluster, exp. task duration

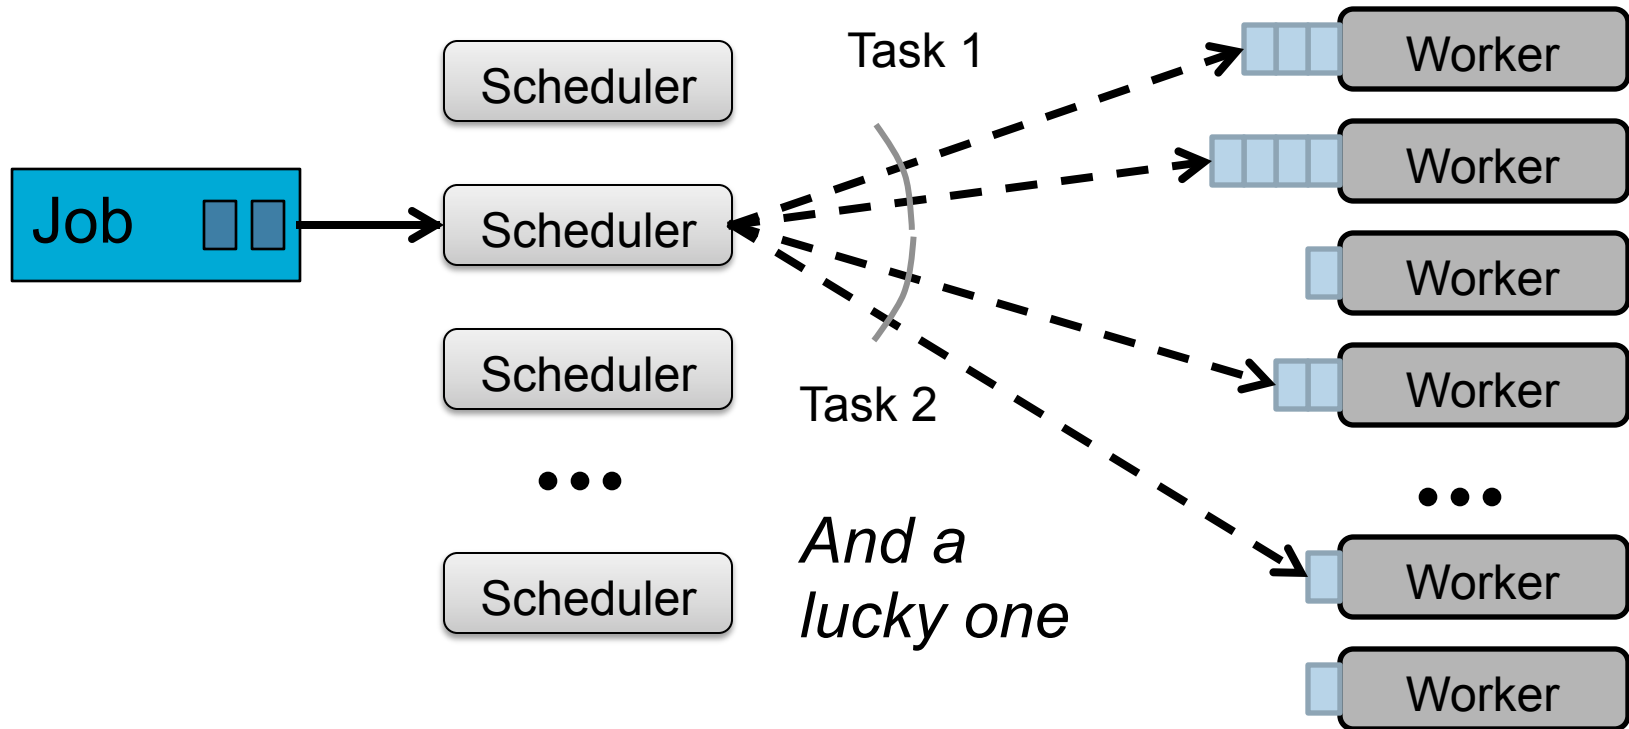
Response time grows with tasks/jobs



70% cluster load

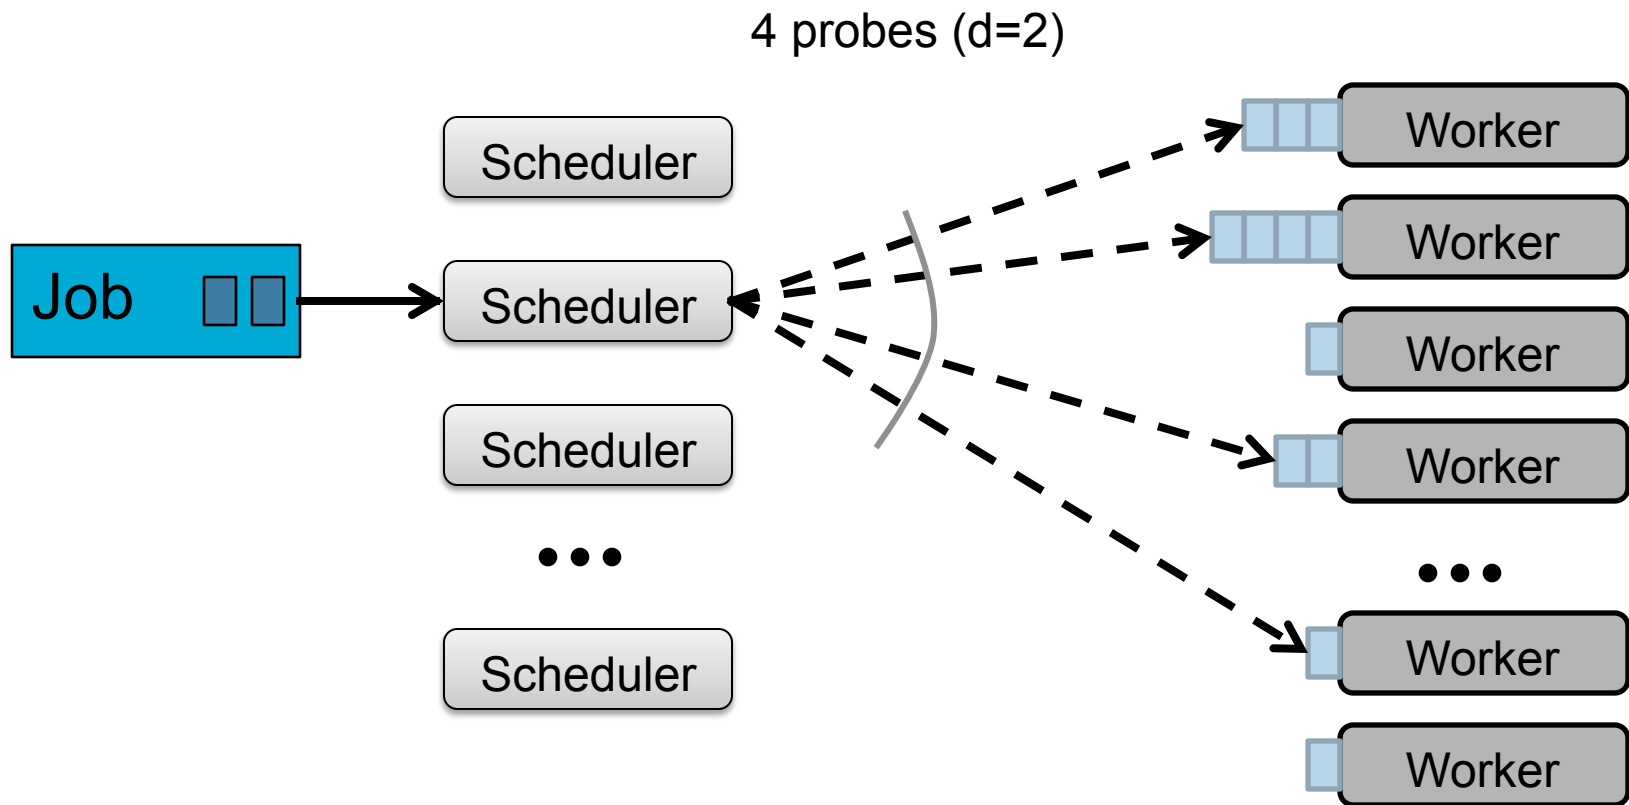
Per-task sampling

An unlucky draw



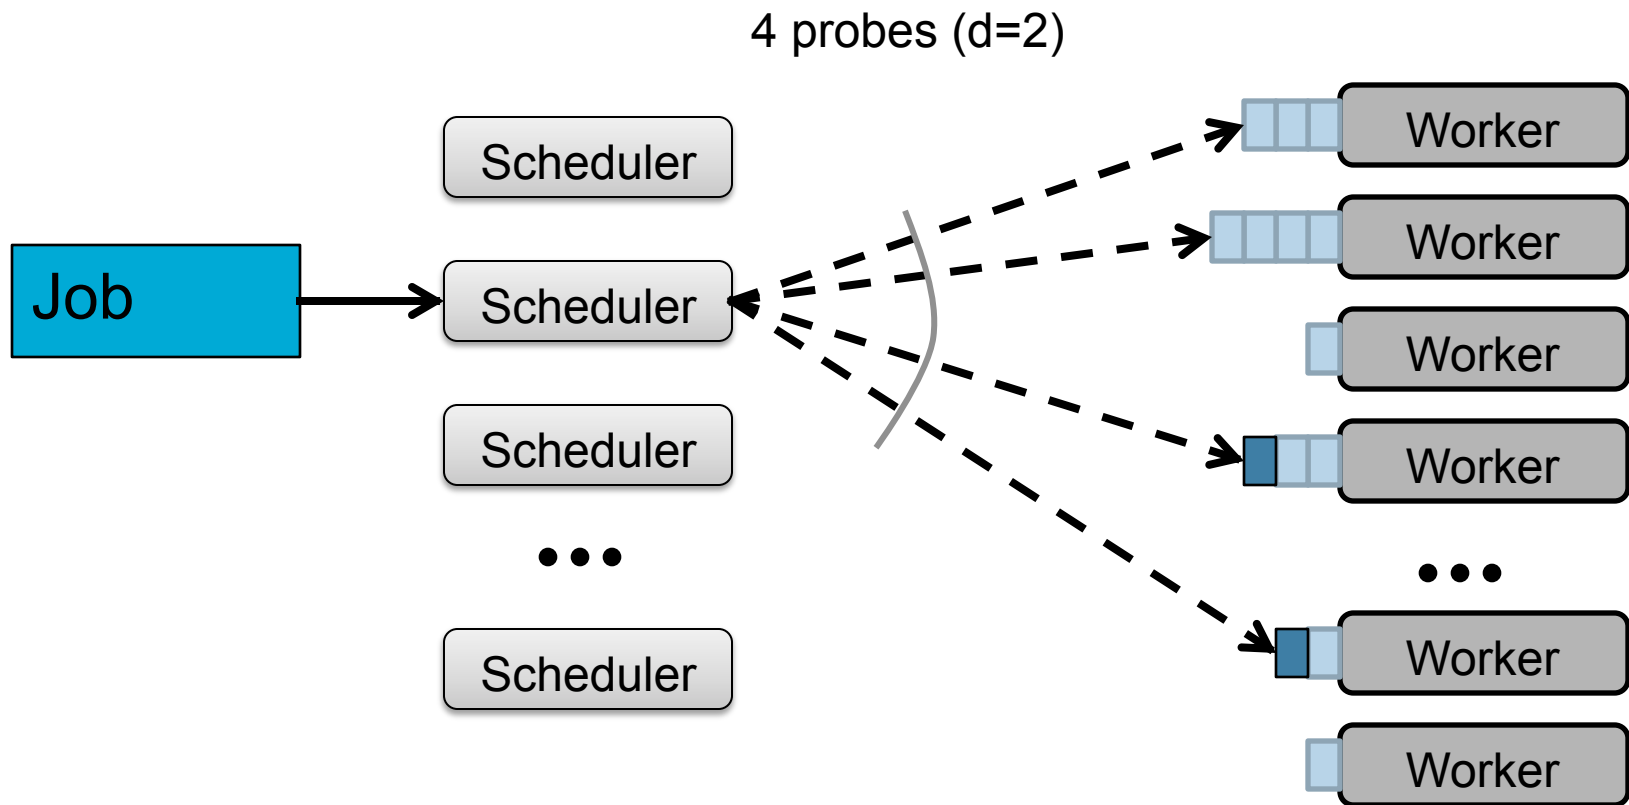
Job is done when all tasks are done ...

Batch sampling




Place m tasks on the least loaded $d*m$ workers

Batch sampling

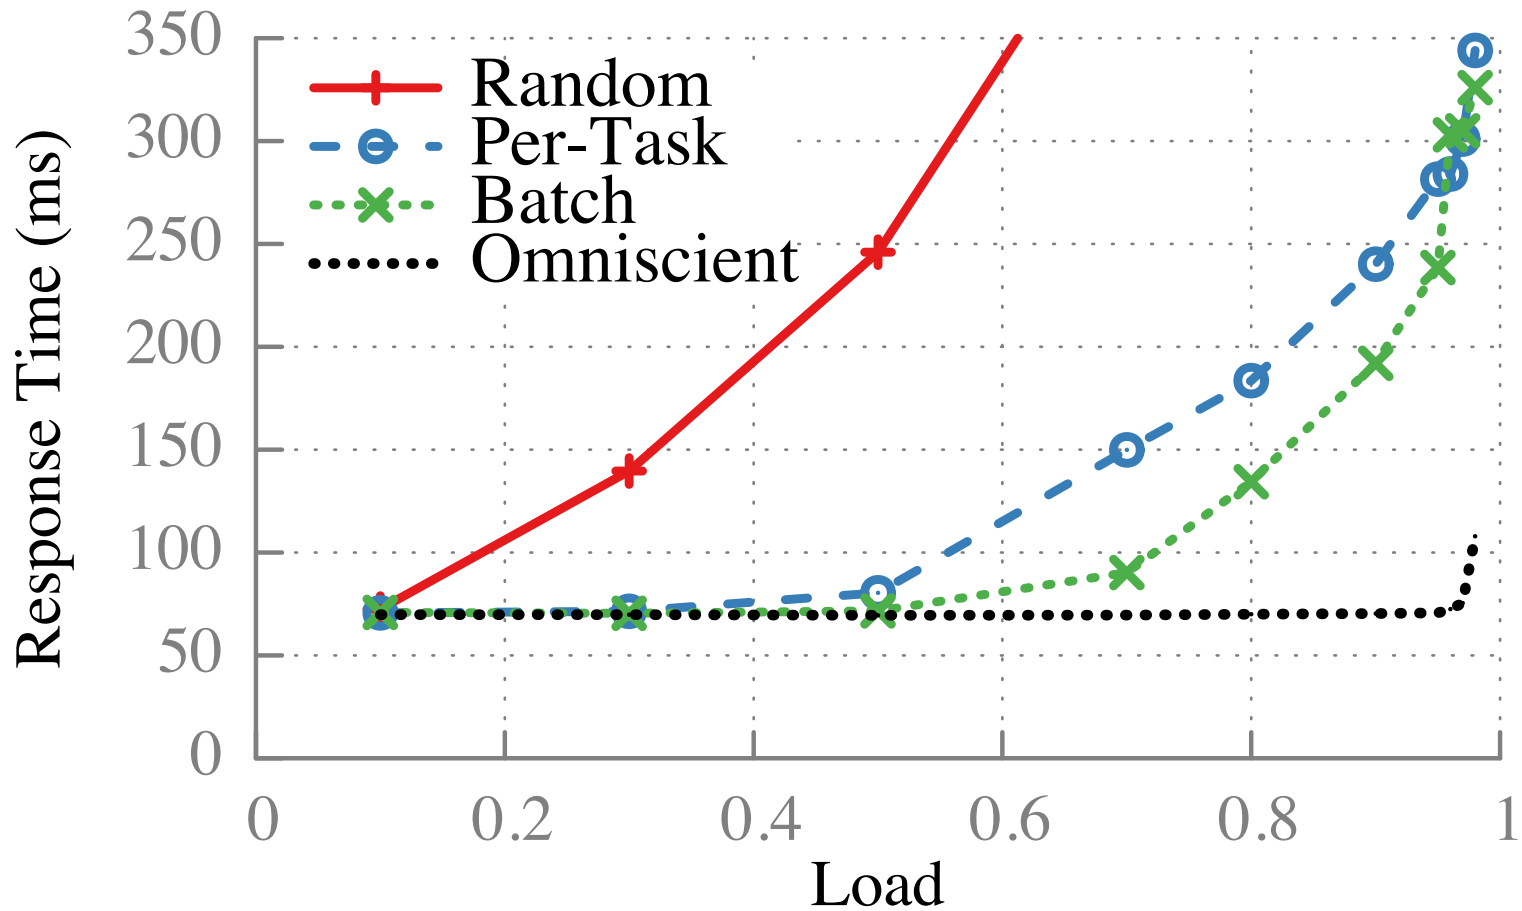


Per-task and Batch



70% cluster load

Batch sampling – better, still 1.92x opt



Late binding

- Sample-based scheduling performs poorly under load
 - Select based on queue length, coarse predictor of wait time
 - Better predictors using estimates of task durations is hard
 - Race condition – multiple schedulers picking the same worker

Late binding


- Late binding
 - Workers put task internal work queue, hold back response
 - When task gets to the front, worker replies RPC
 - Scheduler assigns job's tasks to first m workers
 - ... and sends no-ops to the rest (proactive)
- Cost
 - Idle while sending RPC – a 2% efficiency loss
 - Fraction of time idle while requesting tasks

$$(d * \text{RTT}) / (t + d * \text{RTT})$$

Mean task service time

Mean network round trip time

Late binding benefits



Handling constraints and policies

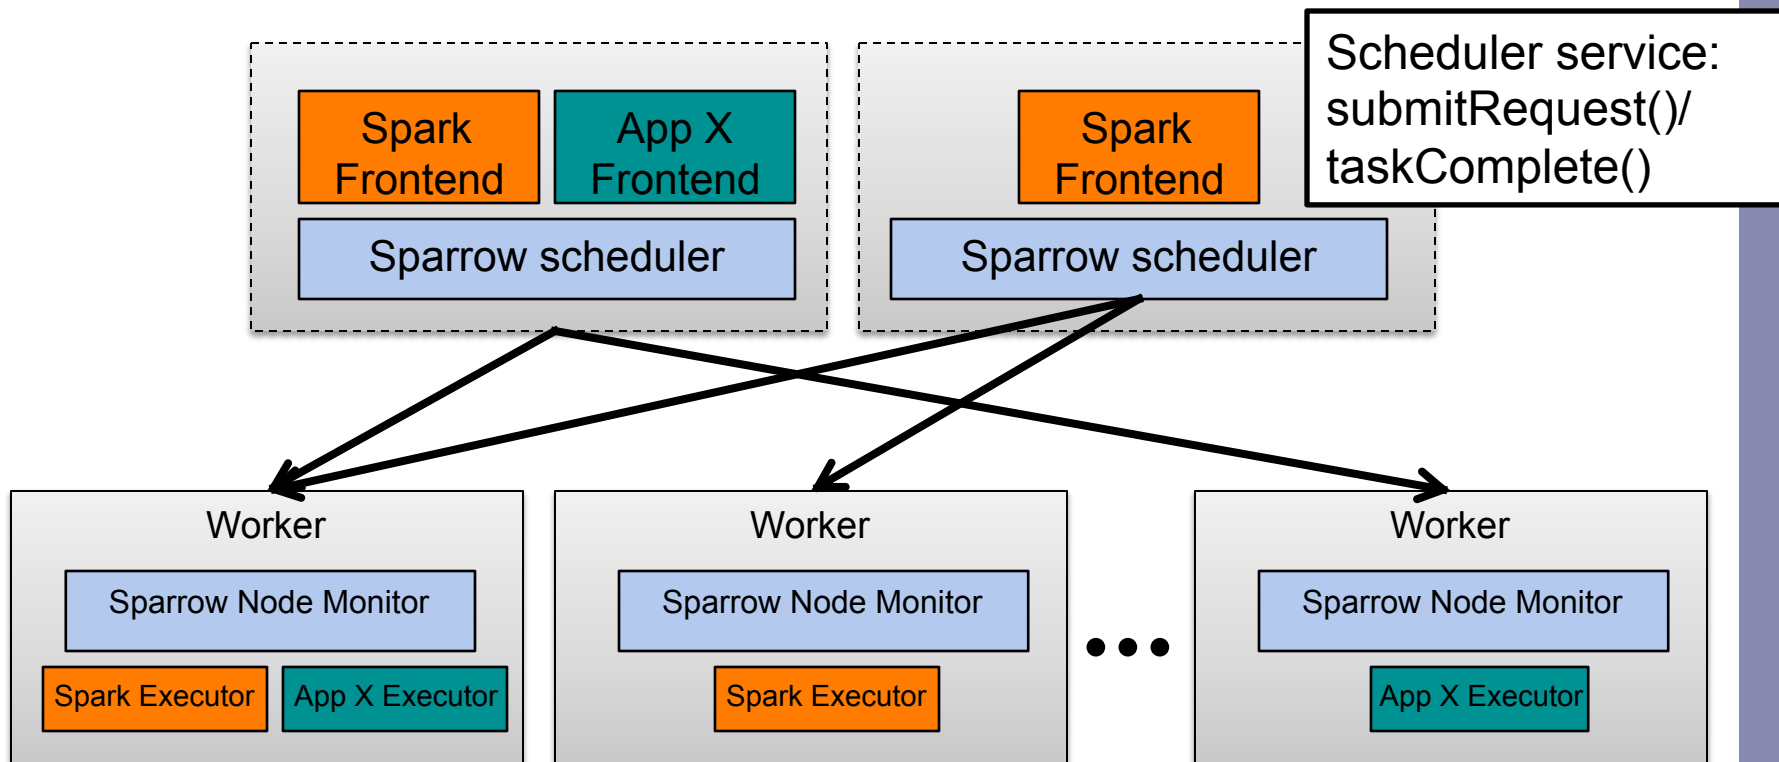
- Job constraints, e.g., *run in workers with GPU*
 - Trivially handled, pick *dm* candidate from the subset
- Per-task constraints, e.g., *run where input is*
 - Use per-task sampling, improved with
 - ... sharing information across tasks when possible
 - ... use late-binding

Handling constraints and policies


- Resource allocation policies, under load
 - Strict priorities through multiple queues on workers
 - FIFO, earliest deadline first, shortest job first, class, ...
 - Trade accuracy for simplicity (no global information but low priority jobs can run first)
 - No preemption
 - Weighted fair sharing
 - One queue per user

Sparrow implementation

- Sparrow works for 1+ concurrent frameworks
 - Front end + executor (long-lived processes responsible for executing task w/o startup overhead)
 - Node monitor federates resources usage bet/ co-located frameworks



Sparrow and Spark's native scheduler



Small 100 16-core EC2, 10 tasks/job,
10 schedulers, 80% load
(synthetic workload)

TPC-H queries: background

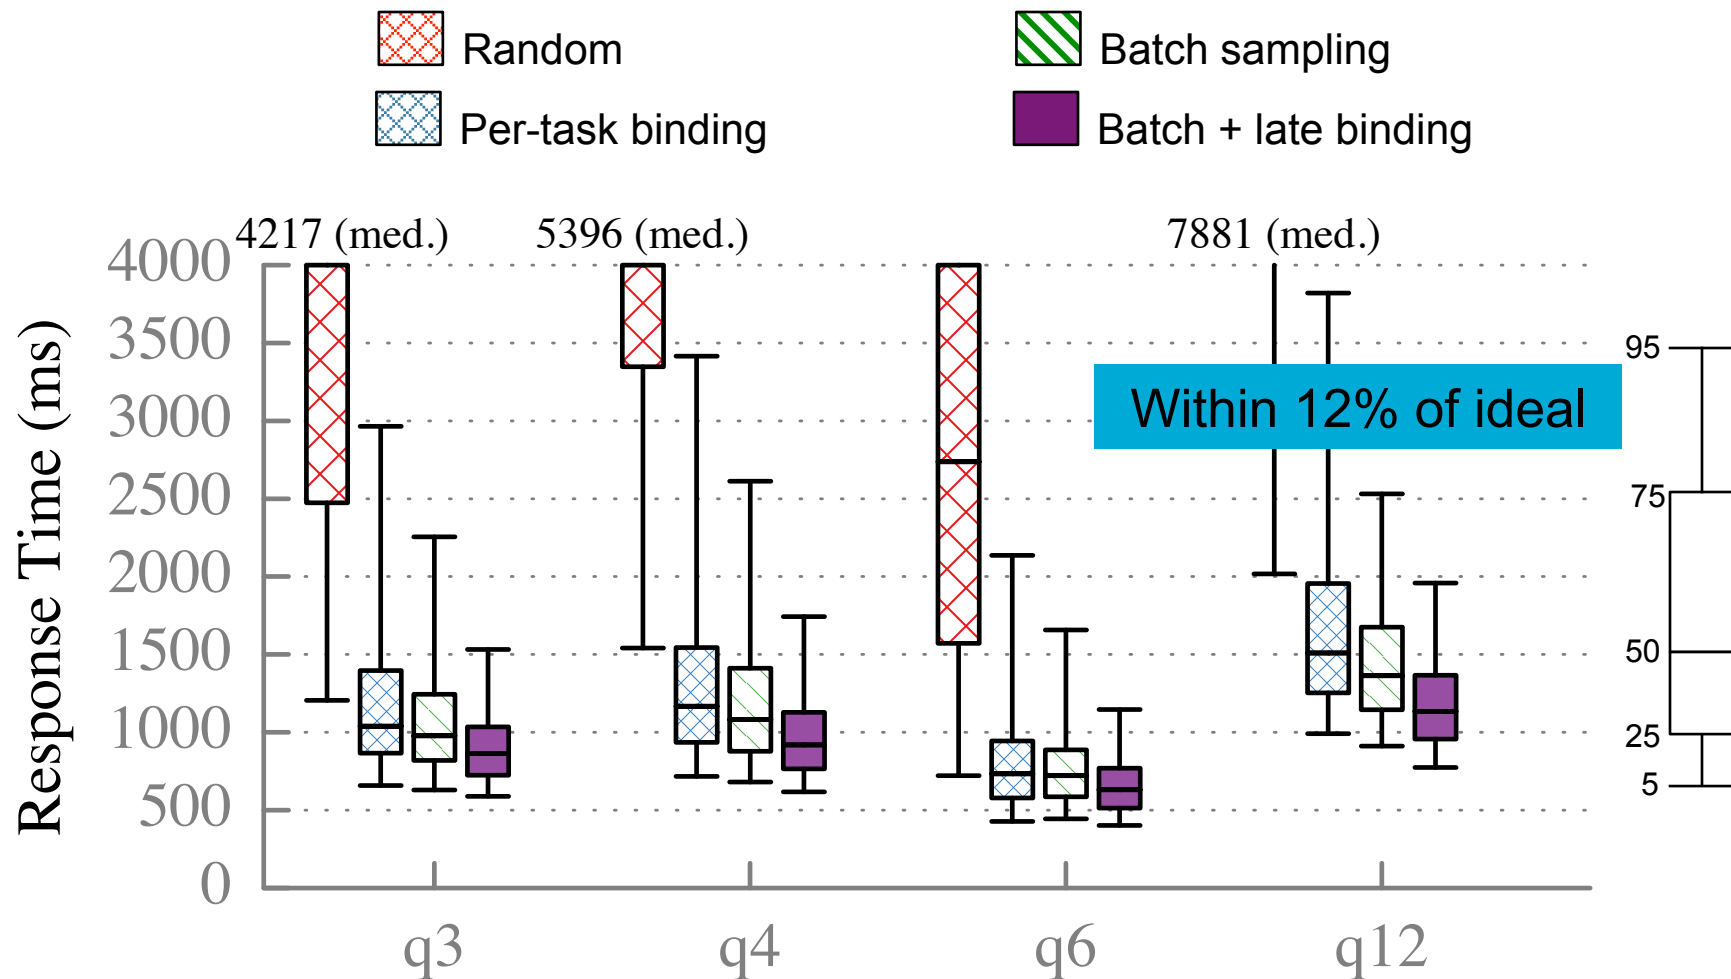
- A common benchmark for analytics workloads; representative of ad-hoc queries on business data
- Shark queries compiled into multiple spark stages
- Each stage triggers a scheduling request using `submitRequest()`
- Task in first stage constrained to machines holding input data
- Stages have different number of tasks, durations and un/constrained queries

Shark: SQL execution engine

Spark: Distributed in-memory analytics framework

Sparrow

TPC-H Queries




100 16-core EC2, 10 tasks/job, 10 schedulers, 80% load


Effect of probe ratio

Too low, can't find lightly loaded machines


Too high, pay the cost of increased messaging



Failure impact



Detect failure 100ms
Failover 5ms
Re-launch queries 15ms



Limitations/future work

- Scheduling policies – can they do better than approximate?
- Inter-job constraints (e.g., *tasks of job A cannot run with those of B*) – hard to do w/o drastic changes
- Gang scheduling – no central point where to do it
- Query-level policies – easy to extend, FIFO
- ...
- *Want to try?* <http://github.com/radlab/sparrow>